

A Study on the Effects of Representation Bias on AI Performance and Methods of Mitigating it

Sarah Abdelazim
Morris Chan
Julie Song

Mentor: Simon Goring



Introduction & Objectives

Data Science Techniques

Data Product

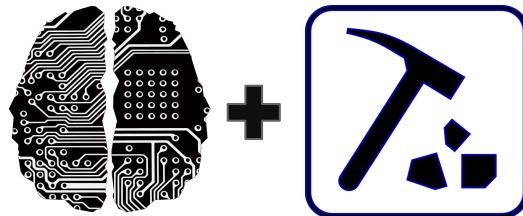
Results

Conclusions and Recommendations

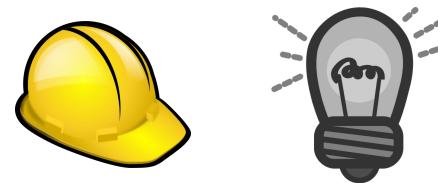


“Mineral exploration is not a new business. But with a stalled rate of discoveries, the global supply chain is faced with dwindling supply and surging demand. With fewer, smaller and more expensive discoveries, the industry is finally doing things differently.”

– <https://goldspot.ca/>



Combines AI with
mineral exploration
projects



Reduce risk, increase
efficiency



Unlock full potential of
historical data

The Dataset

Class 2



Class 3



Class 4

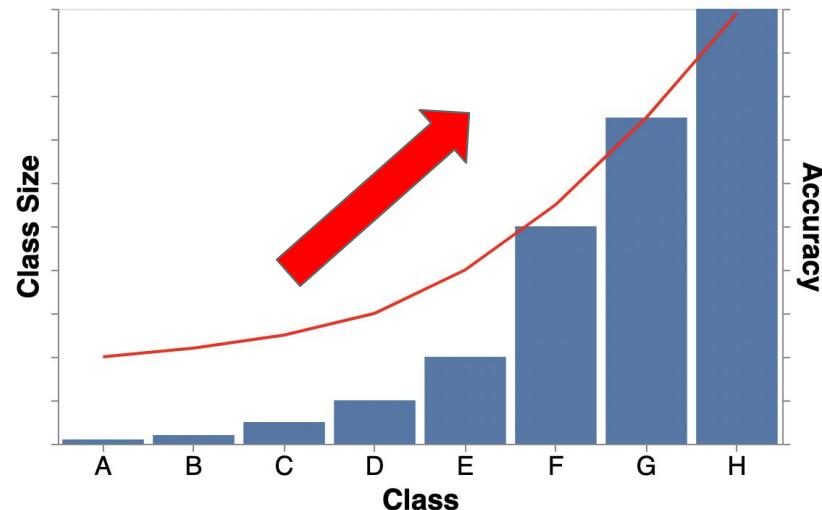
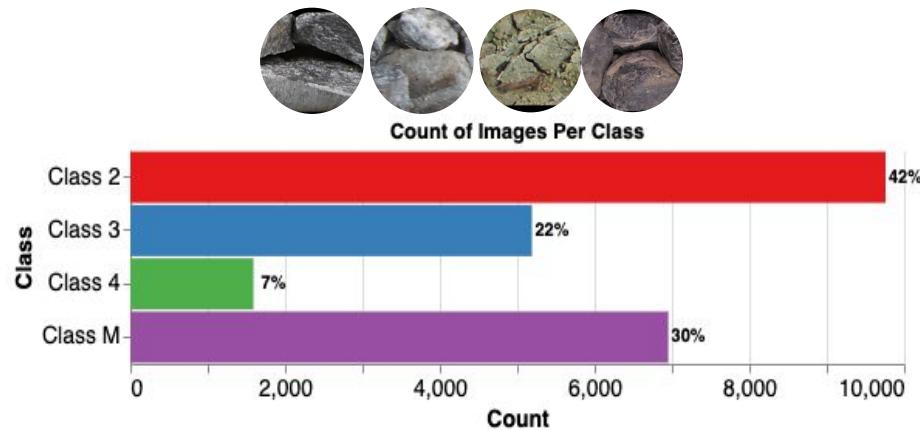


Class M



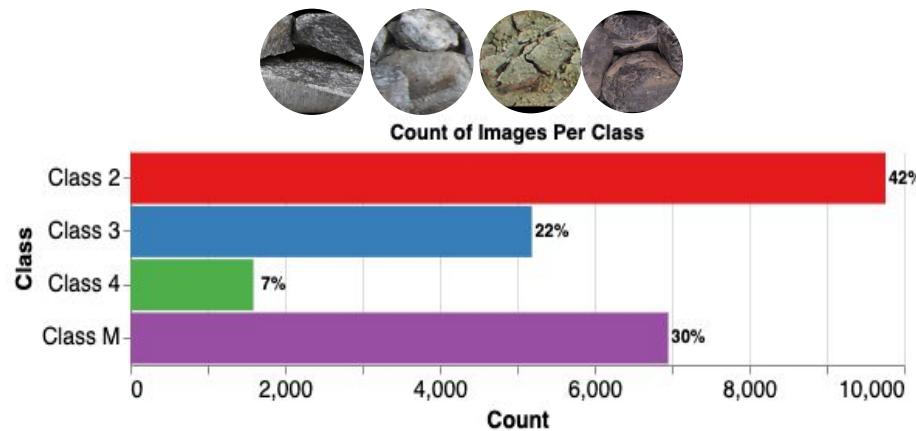
Mining core box ([ref](#))

Data Imbalance



“How can we make accurate predictions
for all 4 classes?”

Data Imbalance

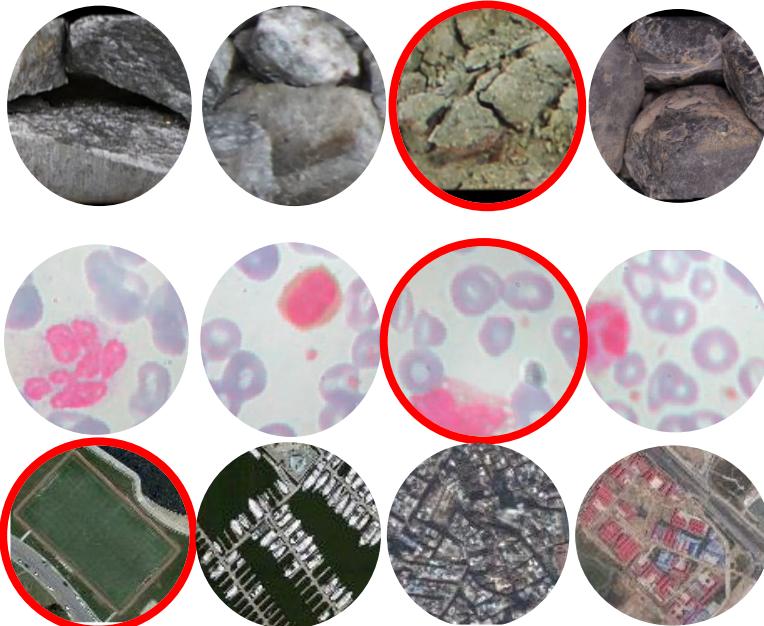


“How can we make accurate predictions for all 4 classes?”

		Predicted Class	
		Safe	Risky
Actual Class	Safe	Good!	Opportunity Costs of New Projects
	Risky	Accidents and Injuries	Good!

The General Problem

Data imbalance is a very general problem.



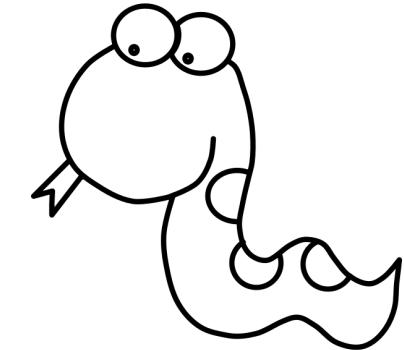
Potential Solutions:

- ★ Data preprocessing
- ★ Model selection
- ★ Loss function implementation
- ★ Optimizer choice

Scientific Objectives



`import debiaser`



- ★ A package that is able to train and compare models, and use different techniques based on the potential solutions to tackle data imbalance
- ★ Success Criteria:
 - Interpretable results
 - Improved performance on imbalance dataset

Introduction & Objectives

Data Science Techniques

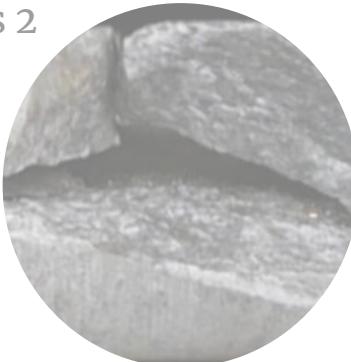
Data Product

Results

Conclusions and Recommendations

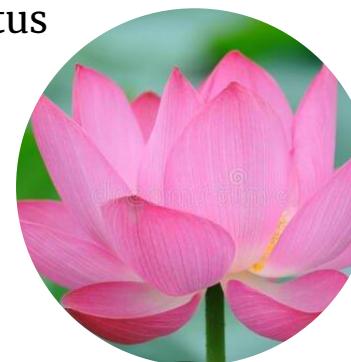
The Dataset

Class 2



Class 3

Lotus



Class M

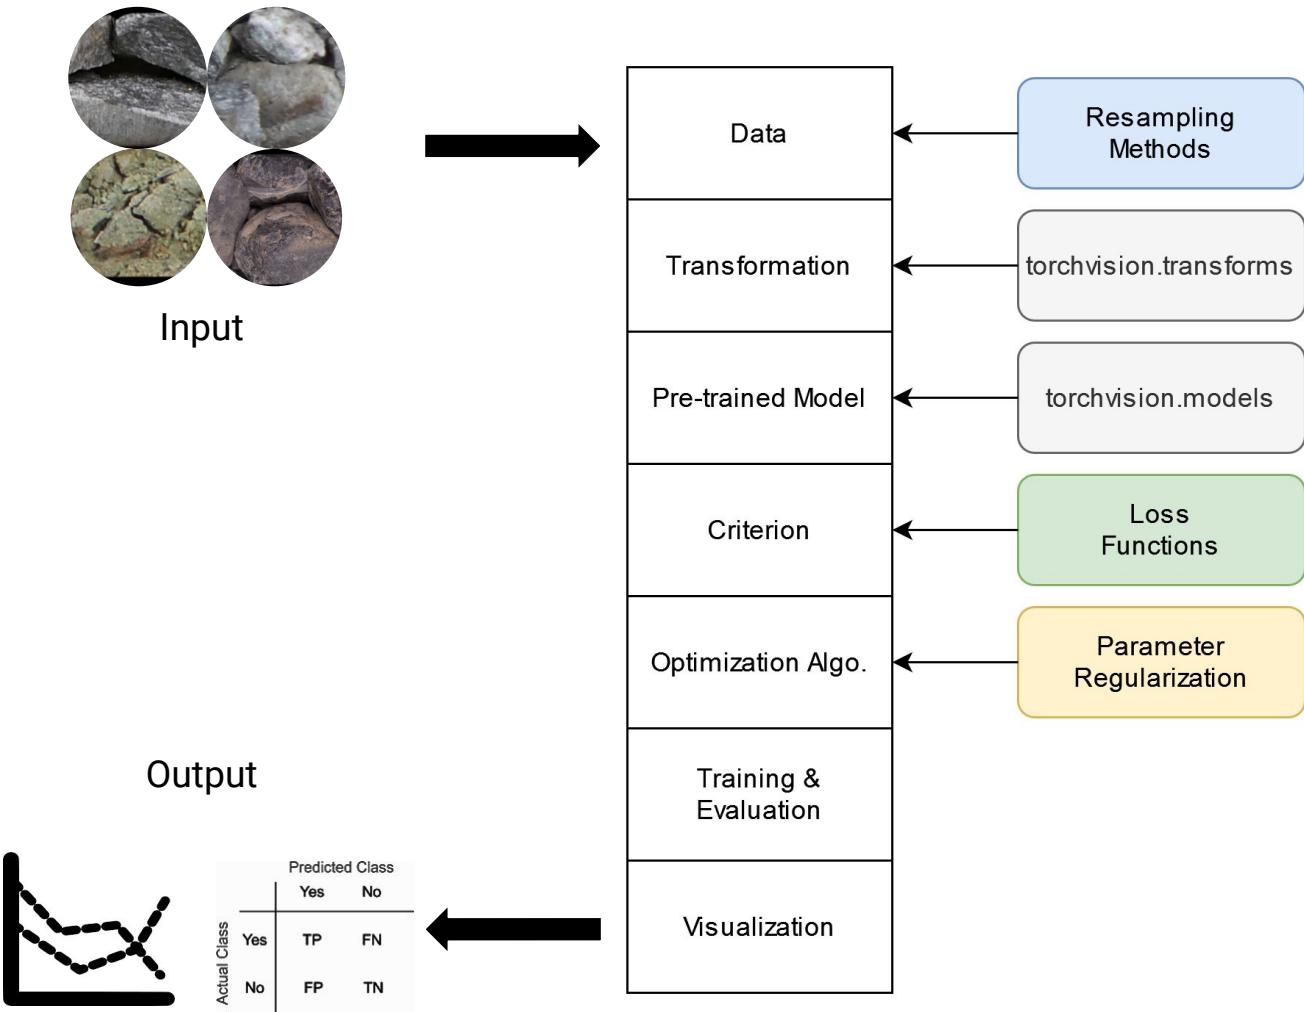
Tulip



Orchid

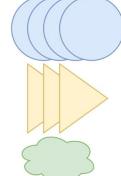
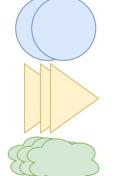
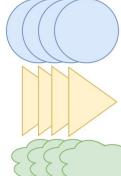
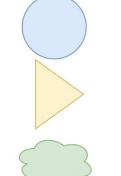
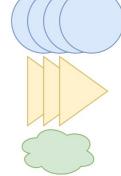
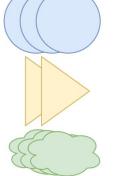
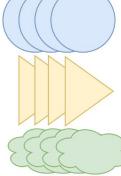
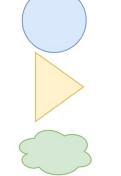
Sunflower

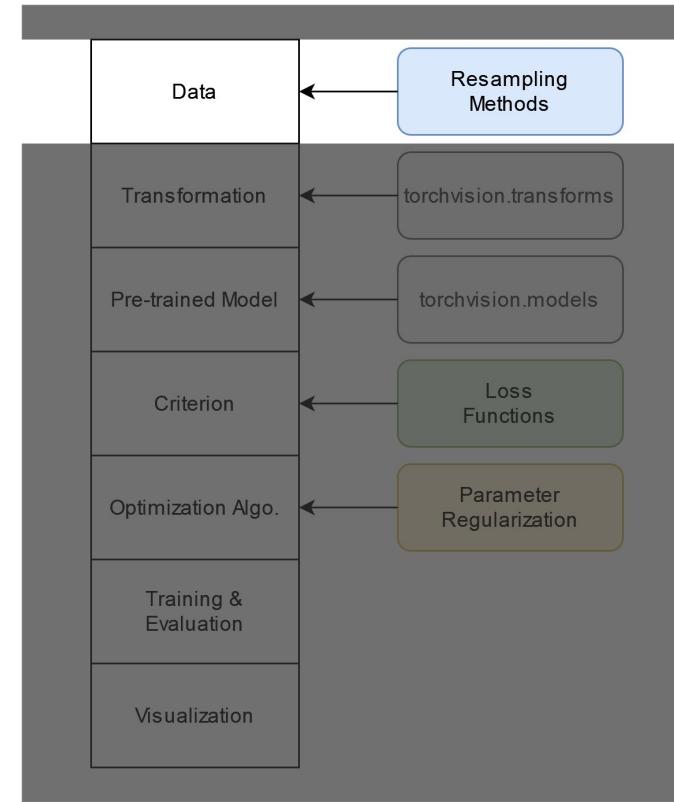
Debiasing Methods Overview



Resampling

```
from debiaser.general import load_image  
  
load_image(  
    ...,  
    sampling_strategy='weighted')
```

Epoch	Original	Weight Random Sampling	Oversampling	Undersampling
1				
2				



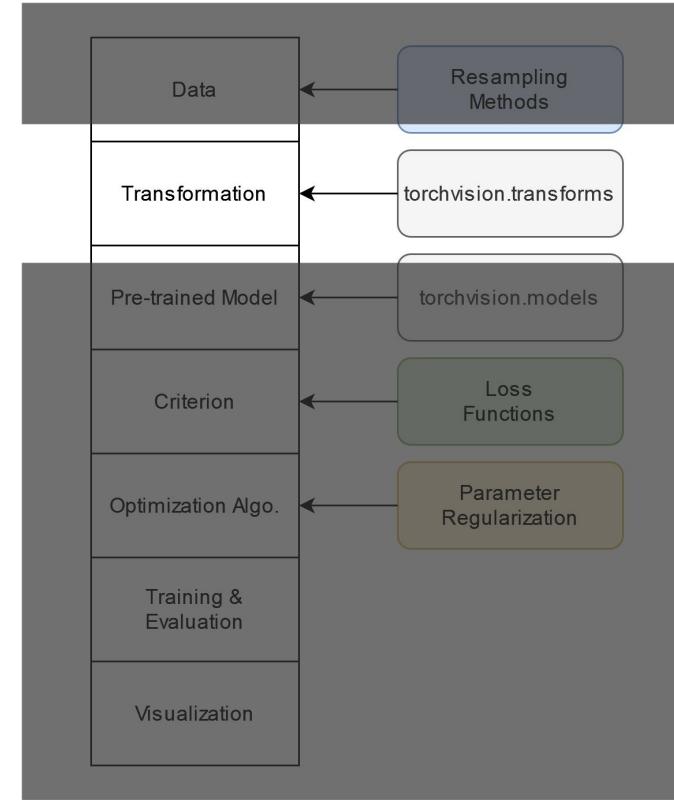
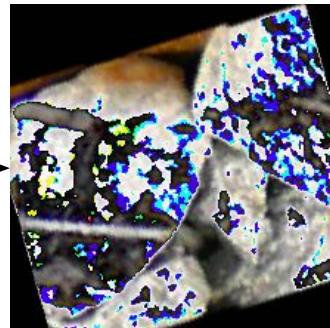
Data Transformations

```
from debiaser.general import transformer  
  
transformer(image_size=224,  
            normalize=(0.5, 0.5),  
            vflip=True,  
            hflip=True,  
            rotate=True,  
            degrees=20,  
            crop=False)
```

Before transformation



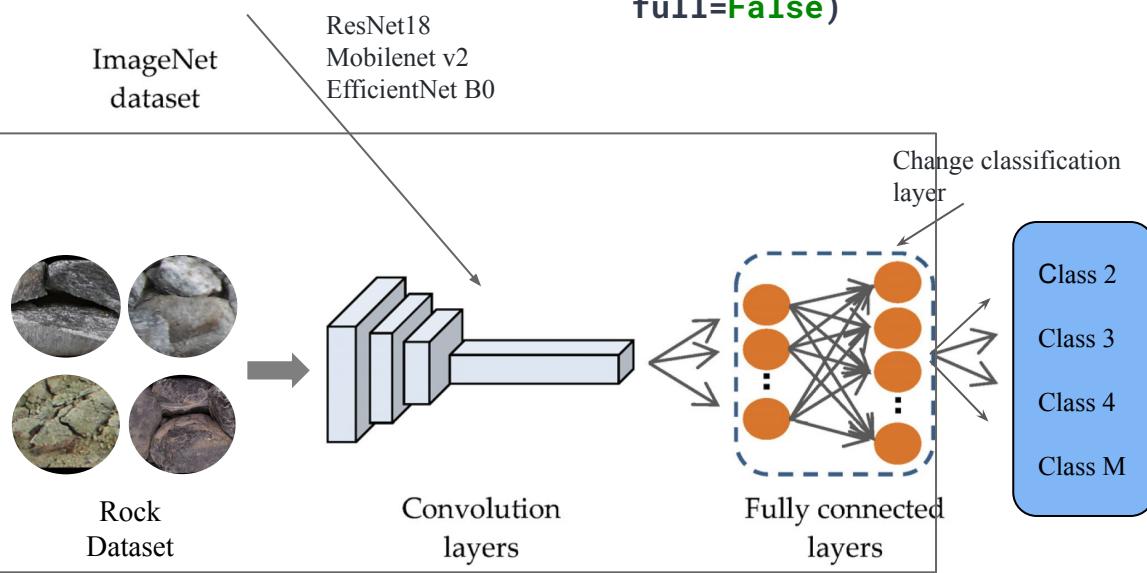
After transformation



Pre-Trained Model

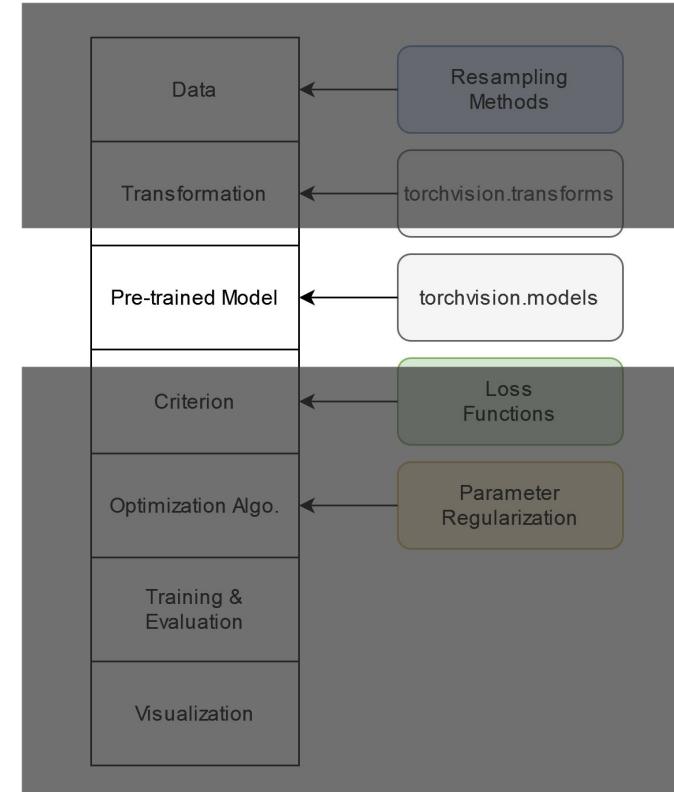


ImageNet
dataset



```
from debiaser.general import load_model

load_model(model_id='resnet18',
           n_class=4,
           pretrained=True,
           full=False)
```



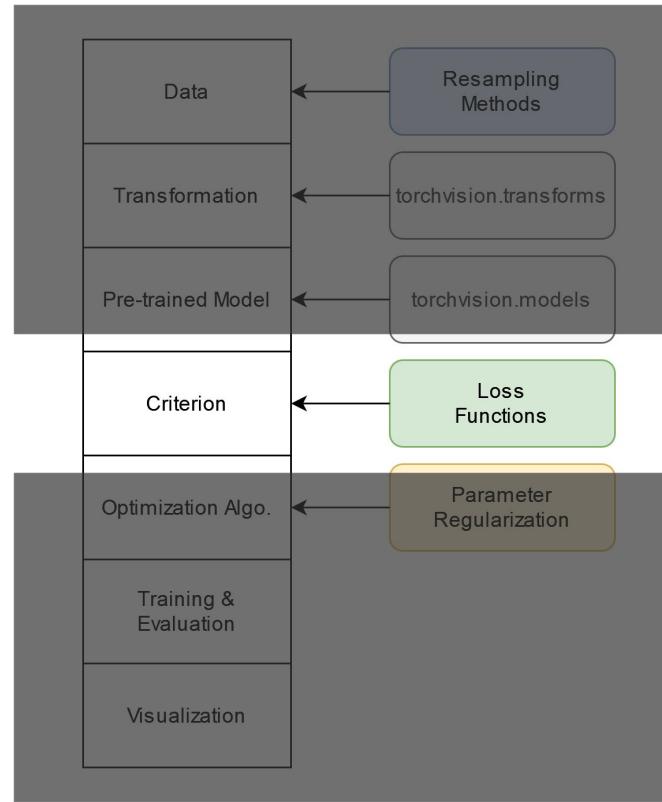
Loss Functions – Context

Problem:

- A *naive* model focuses on learning the characteristics of the majority class while, neglecting the examples from the minority → majority class contributes more to loss function

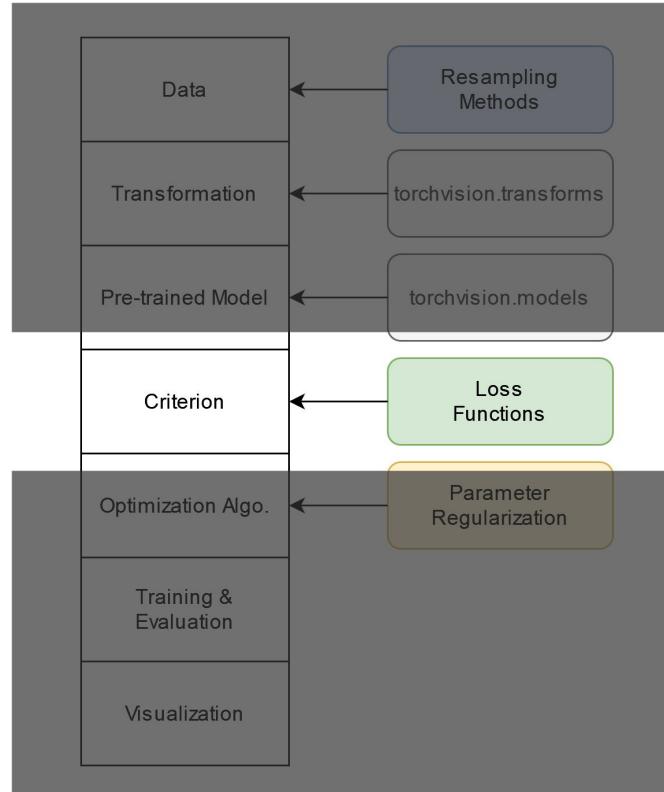
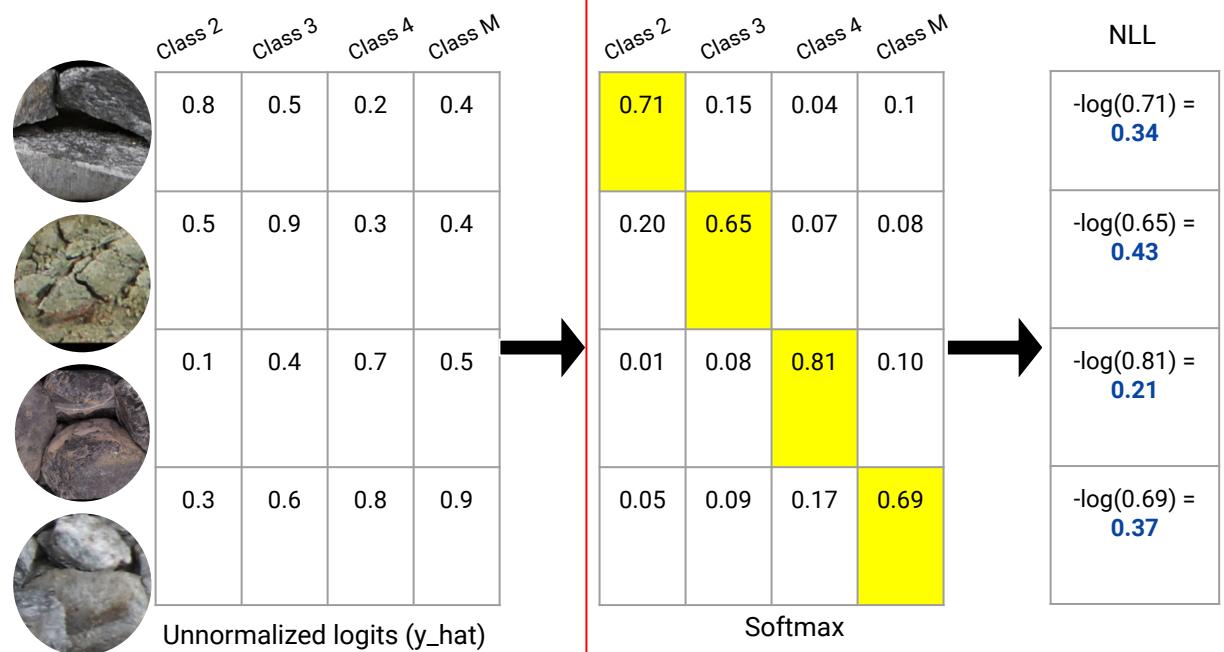
Solution:

- The key to addressing imbalanced datasets is:
 - Class-balanced loss via a weighting scheme (4 approaches)
 - make the loss function more sensitive to errors from the minority class (e.g. Mean False Error)



Loss Functions – Baseline Loss

```
from debiaser.general import load_loss ==> CrossEntropyLoss()  
load_loss(loss_family='baseline', ...)
```



Loss Functions – Class-Balanced Loss

- 1) Pick a weighing scheme:

Inverse of Square Root of Number of Samples (ISNS) Inverse of Effective Number of Samples (ENS)

Inverse Number of Samples (INS)

$$w_{n,c} = \frac{1}{\text{Number of Samples in Class } c}$$

ISNS

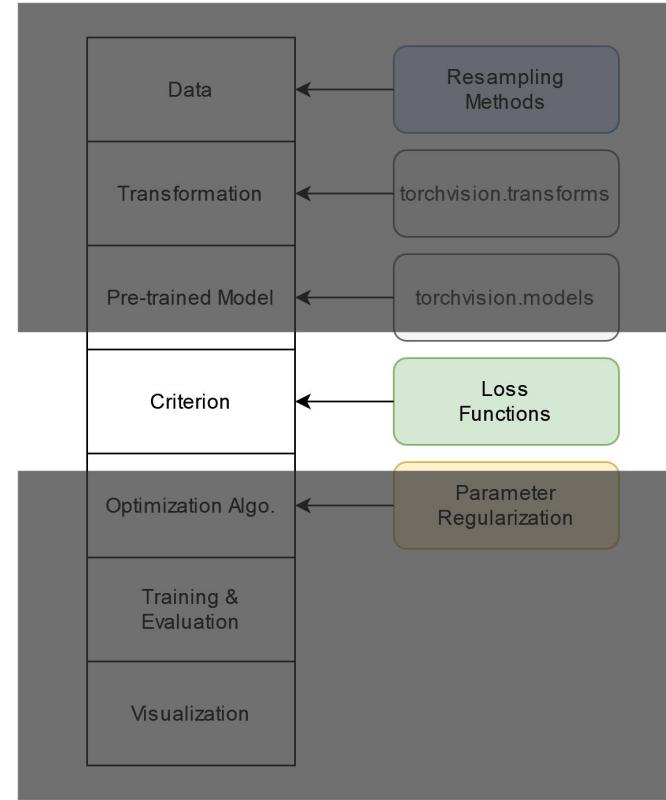
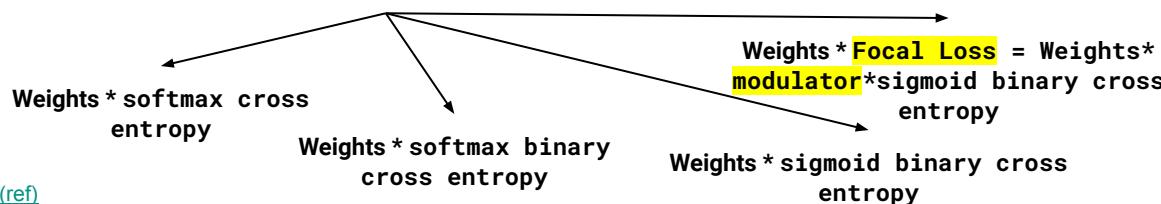
$$w_{n,c} = \frac{1}{\sqrt{2 \cdot \text{Number of Samples in Class } c}}$$

ENS

$$E_{n_c} = \frac{1 - \beta^{n_c}}{1 - \beta} \quad \beta \in [0, 1]$$
$$w_{n,c} = \frac{1}{E_{n_c}}$$

```
load_loss(loss_family='CB', loss_type=..., beta=0.999, gamma=2)
```

- 2) Use with the appropriate loss function:



Loss Functions – Mean (Squared) False Error

```
load_loss(loss_family='MFE', loss_type='MSFE')
```

- Mean Squared Error (MSE):

$$l = \frac{1}{M} \sum_i \sum_n \frac{1}{2} (d_n^{(i)} - y_n^{(i)})^2$$

- Mean false error (MFE): A derivative of MSE

$$FPE = \frac{1}{N} \sum_{i=1}^N \sum_n \frac{1}{2} (d_n^{(i)} - y_n^{(i)})^2$$

$$FNE = \frac{1}{P} \sum_{i=1}^P \sum_n \frac{1}{2} (d_n^{(i)} - y_n^{(i)})^2$$

$$l' = FPE + FNE$$

- Mean squared false error (MSFE):

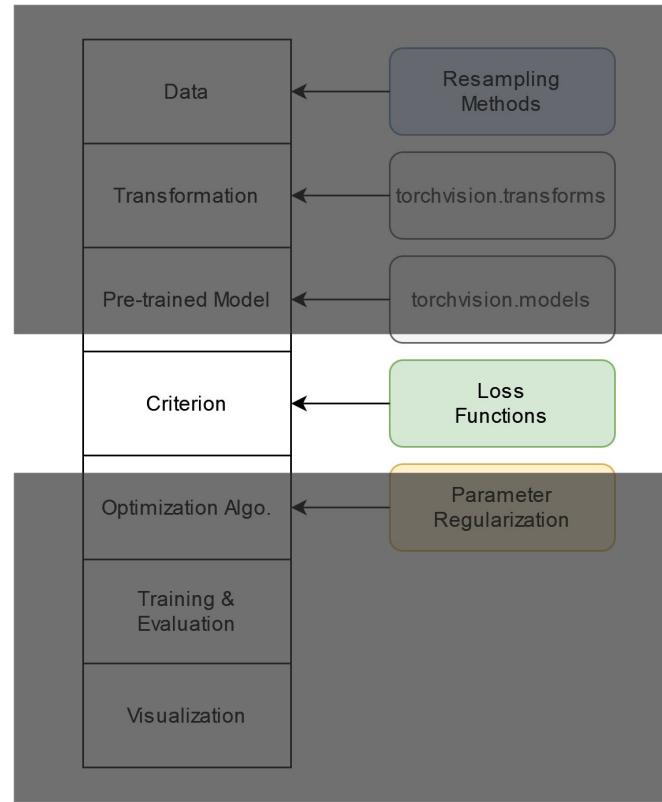
$$l'' = FPE^2 + FNE^2$$

N = sample size of negative class
which is the majority class

P = sample size of positive class
which is the minority class

d_n = the label vector

y_n = the prediction vector which is
the output of sigmoid or softmax
layer



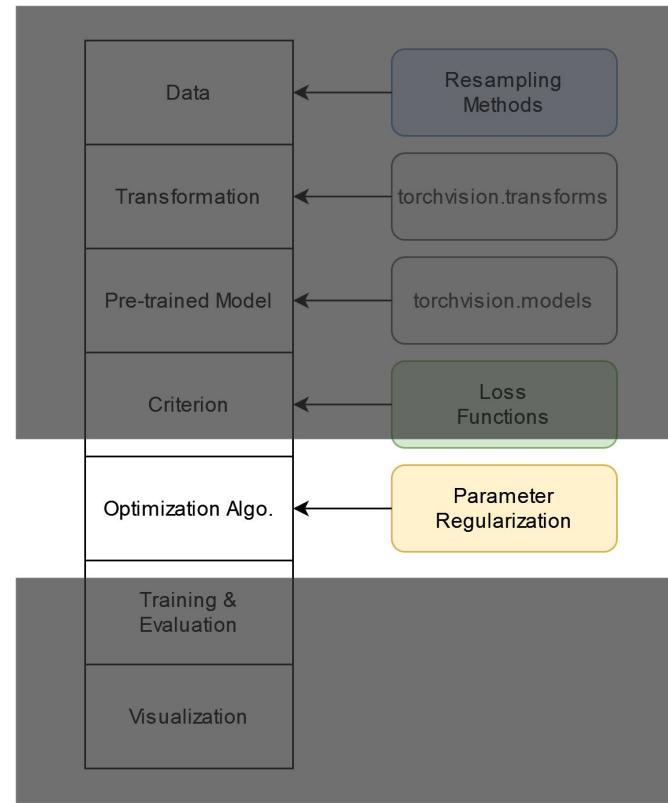
Parameter Regularization - Context

A naive model has large weights for common classes

- Solution: Try to balance the weights across classes for better predictions on the minority class!

```
from debiaser.general import load_optimizer

load_optimizer(dict_model,
               lr=1e-5,
               optimizer_name = 'Adam',
               strategy = None)
```

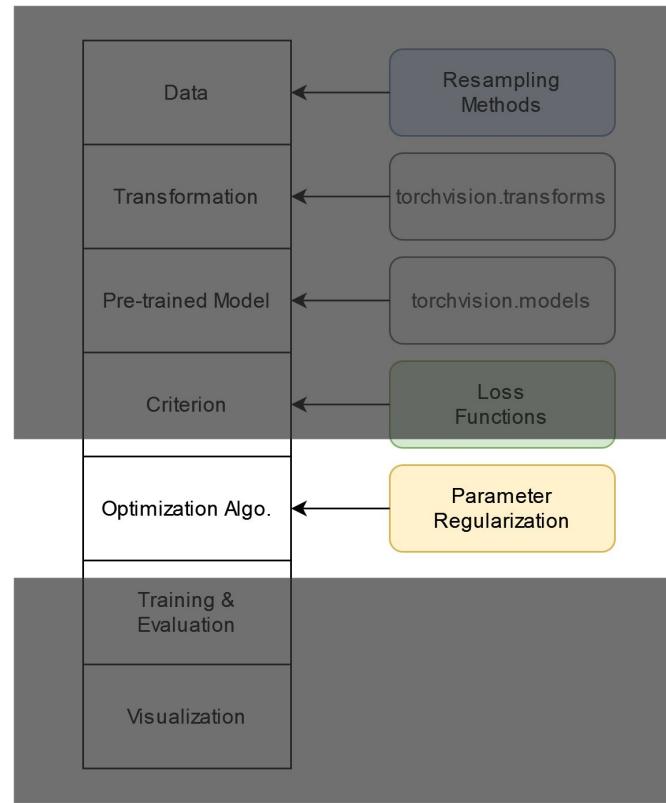


Parameter Regularization - Weight Decay

Weight decay penalizes larger weights more severely than smaller weights

- More balanced weight growth
- Decreases model complexity
- Improves generalization

```
load_optimizer(dict_model,  
              lr=1e-5,  
              optimizer_name='Adam',  
              strategy='wei_dec',  
              weight_decay=0.005)
```

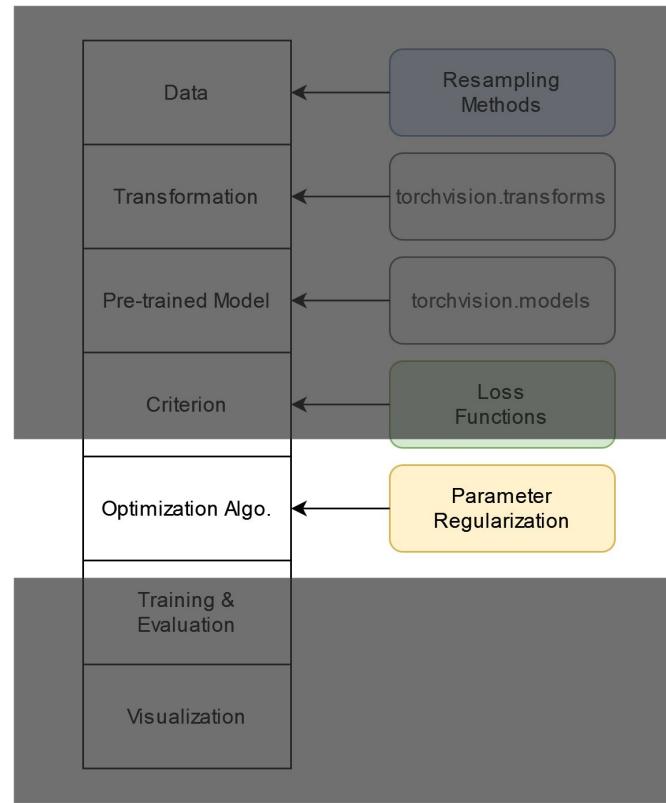


Parameter Regularization - L2-Norm

L2- normalization places a strict restriction on the weights.

- Improves minority class performance
- Sacrifices overall model performance

```
load_optimizer(dict_model,
               lr=1e-5,
               optimizer_name='Adam',
               strategy='l2_norm')
```

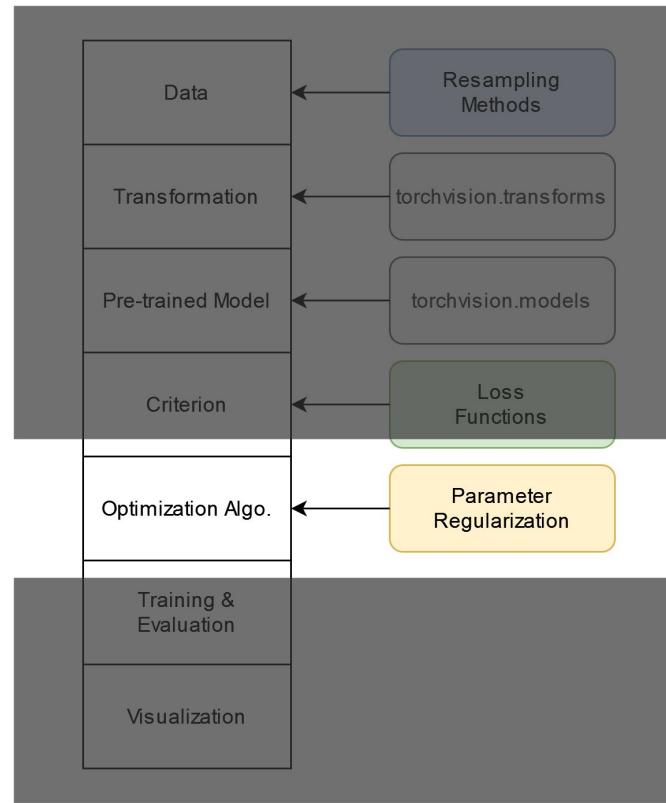


Parameter Regularization - MaxNorm

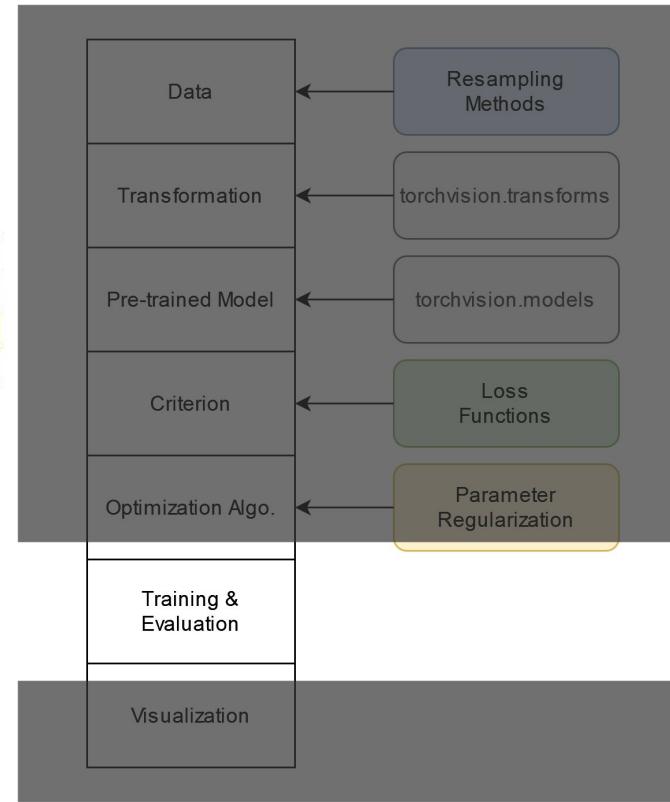
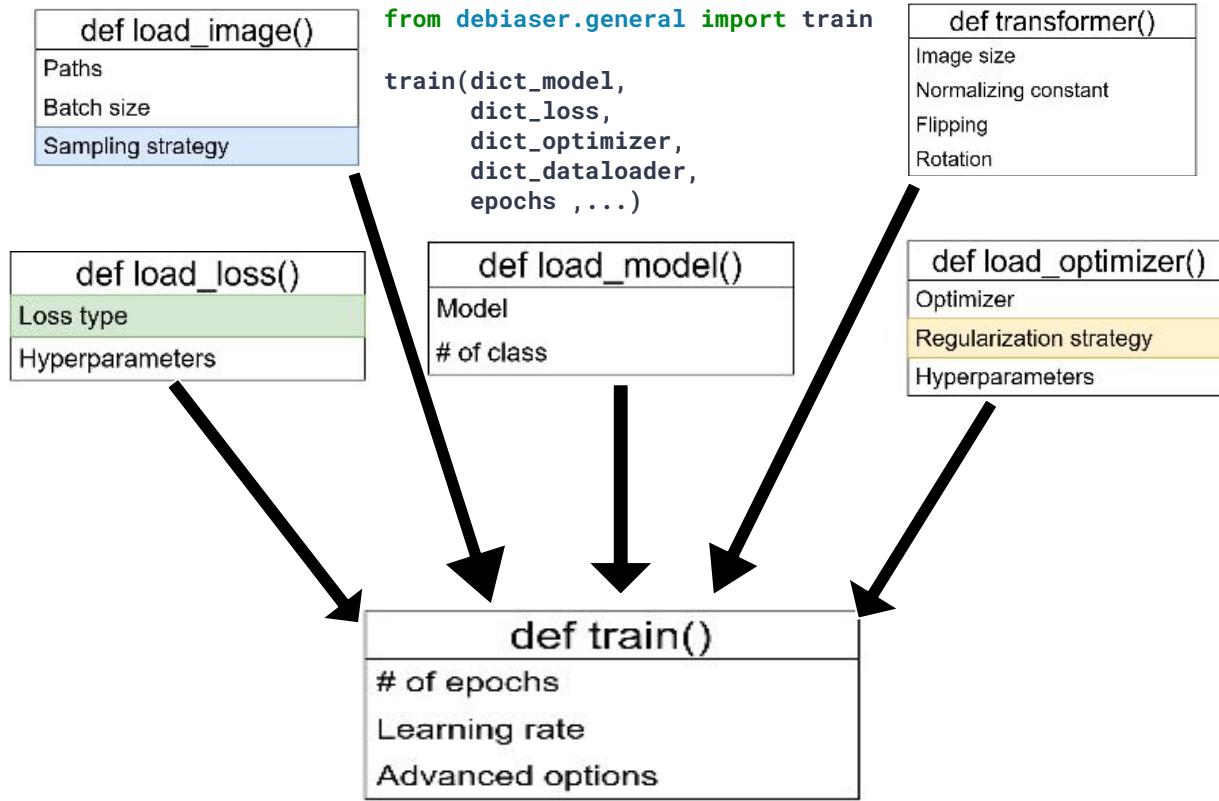
MaxNorm encourages smaller weight growth and places a limit on weight norms (δ)

- Weights are stable even at high learning rates because they are always bounded.

```
load_optimizer(dict_model,
               lr=1e-5,
               optimizer_name='Adam',
               strategy='maxnorm',
               thresh=0.8)
```



Training



Introduction & Objectives

Data Science Techniques

Data Product

Results

Conclusions and Recommendations

User Persona



- Highly technical
 - Familiar with Python and programming
-
- **Generalized!**
 - **Modify easily!**
 - **JSON formats!**

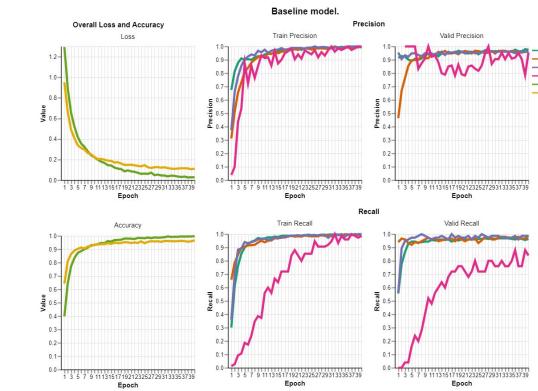
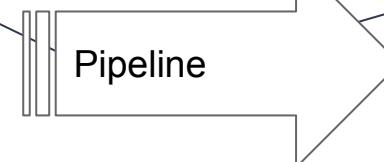
Basic functions

```
load_image    transformer    load_model    load_loss    load_optimizer    train
```

{

```
'batch_size': 32,  
'model': 'resnet18',  
'loss_type': 'focal',  
'beta': 0.9999,
```

```
...
```



Json File Schema

- Json
 - **note** { 'note': 'baseline-resnet',
 - **date** 'date': '20230609-122303,
 - **results** }
 - **y** ...
 - **var_dict**

Json File Schema

- Json
 - note
 - date
 - **results**
 - Training and validation set
 - Loss and accuracy
 - Precision and recall for each class
 - y
 - var_dict
- ```
{
 'results': {
 'train_loss': [...],
 'train_accuracy': [...],
 ...
 },
 ...
}
```

# Json File Schema

- Json

- note
- date
- results
- **y**
  - All target labels & predicted labels
  - Both training and validation set
  - All epochs
- var\_dict

```
{
 'y': {
 'train': {
 'true': {
 0: [1, 0, 1, ...],
 1: [0, 0, 0, ...],
 ...
 },
 'pred' : {...},
 },
 'valid' :{...}
 }
}
```

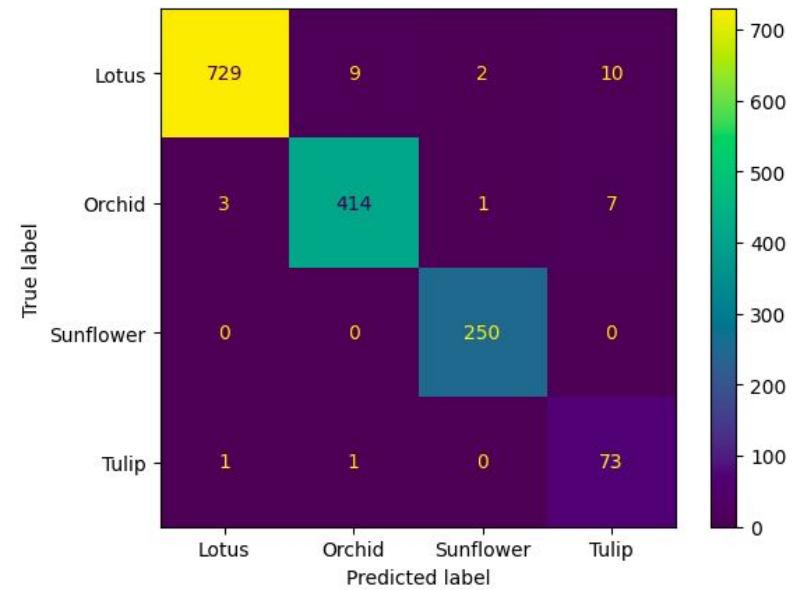
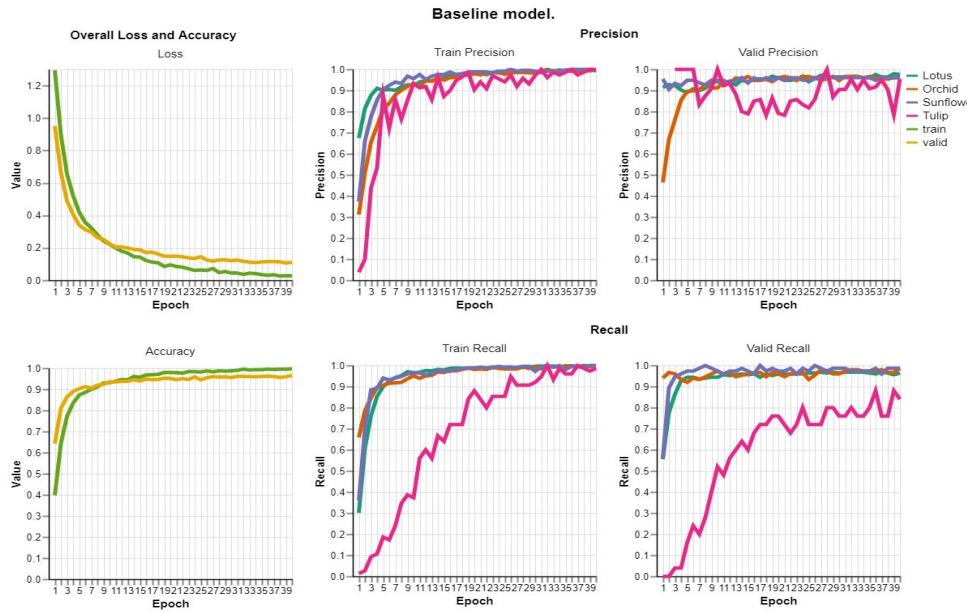
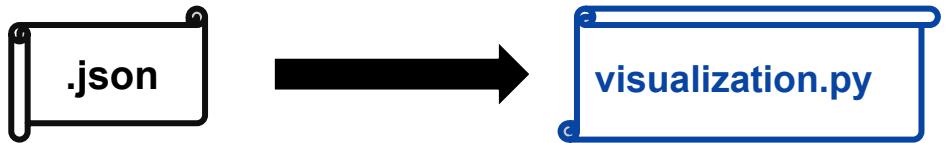
# Json File Schema

- Json

- note
- date
- results
- y
- **var\_dict**
  - All parameters used
  - Identifiers for trials
  - Useful for comparing trials

```
{
 'var_dict': {
 'batch_size': 32,
 'sampling_strategy': None,
 'model_name': 'resnet18',
 'loss_family': 'CB',
 'loss_type': 'focal',
 'beta': 0.999,
 'gamma': 0.5,
 'optimizer': 'Adam',
 'learning_rate': 1e-5,
 'epochs': 50
 }
}
```

# Visualizing One JSON file

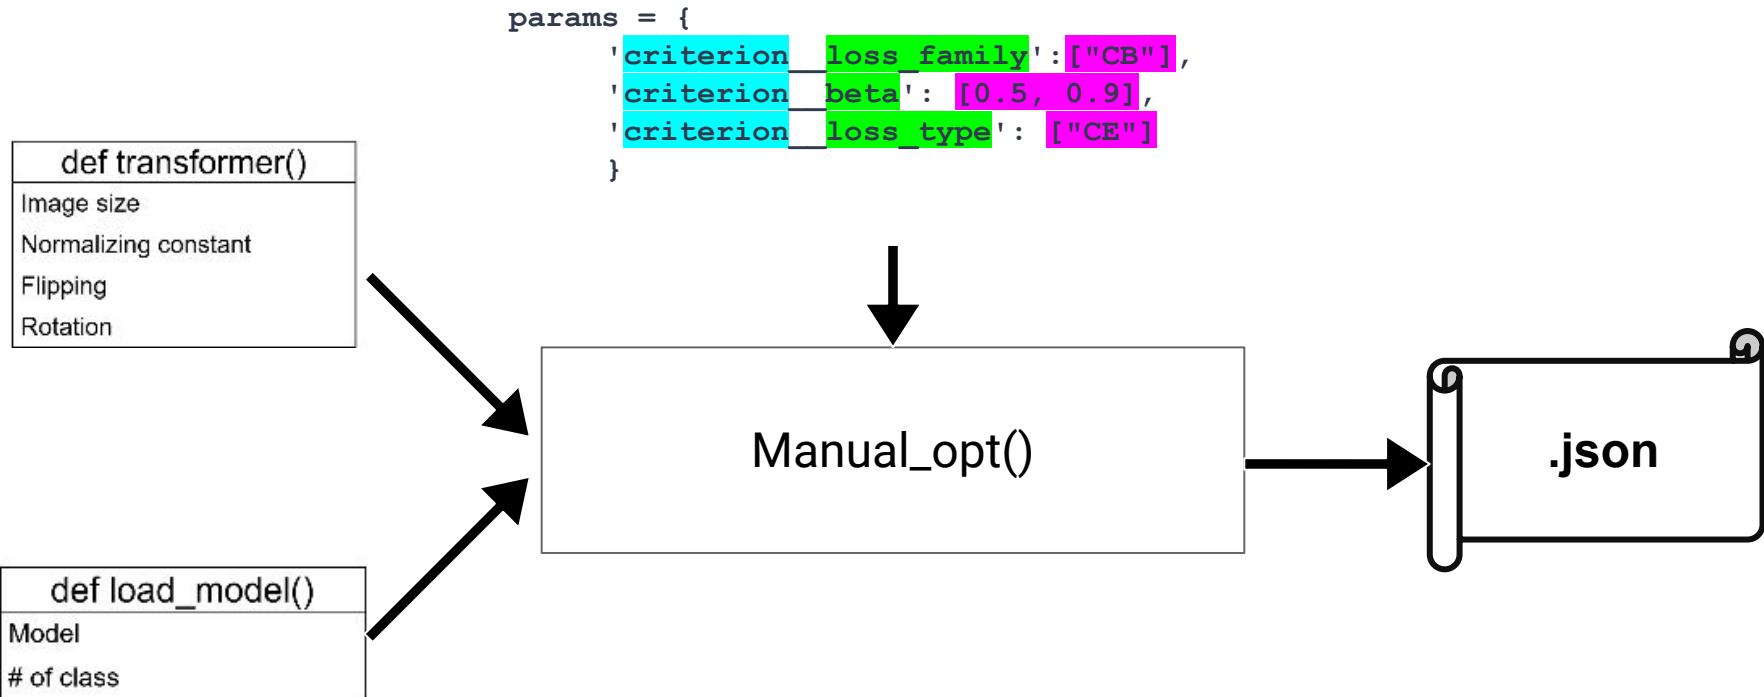


# Hyperparameter Optimization

GridSearch with K-fold cross validation for each combination of parameters

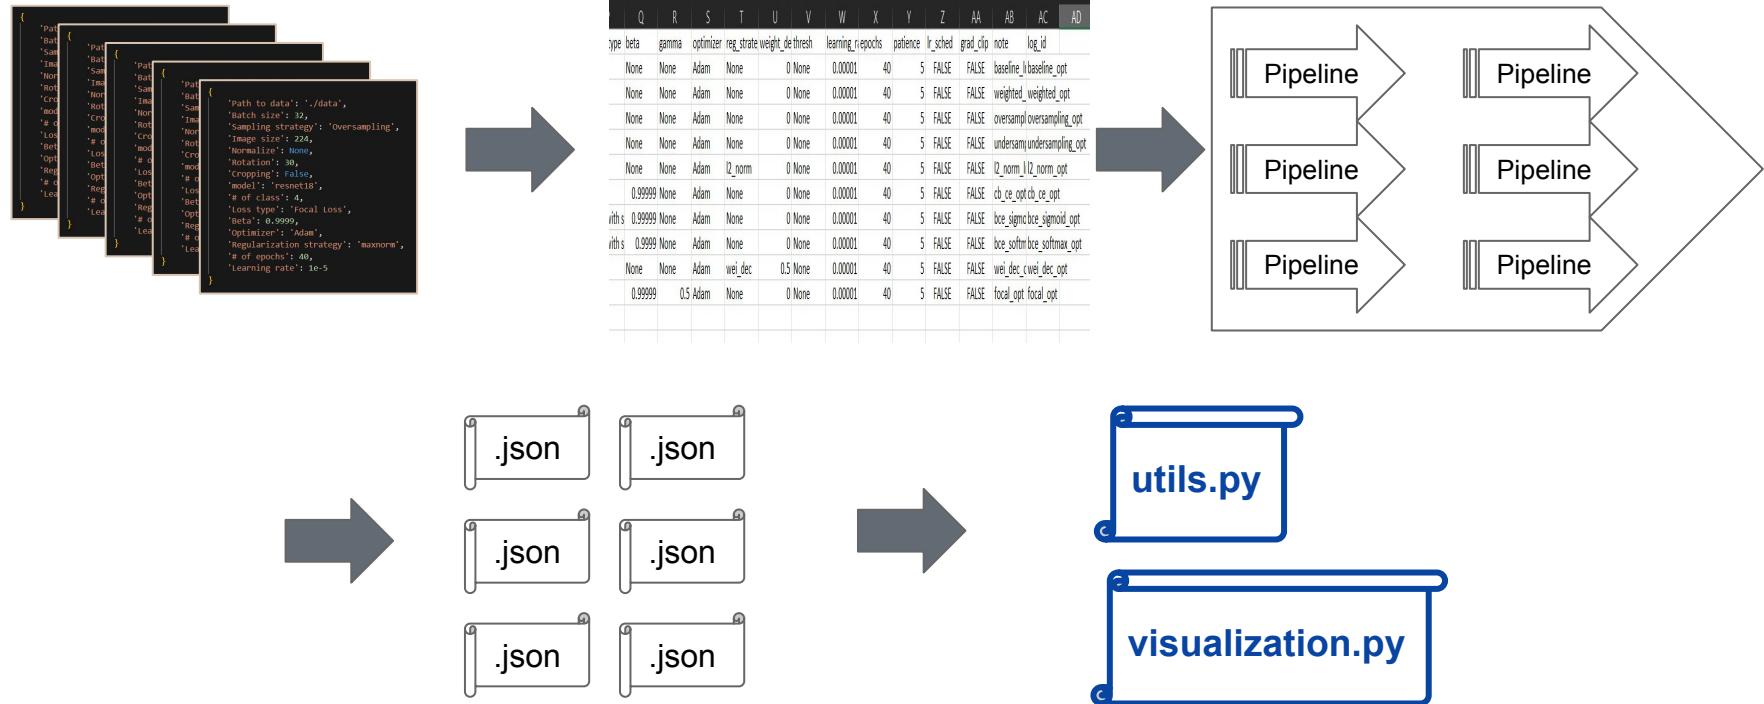
| <b>Optimizer</b>                                                                | <b>Loss Function</b>                              | <b>Train</b>                                                                                 |
|---------------------------------------------------------------------------------|---------------------------------------------------|----------------------------------------------------------------------------------------------|
| Learning rate<br>Optimizer<br>Regularization strategy and associated parameters | Loss family<br>Loss type<br>Associated parameters | Number of epochs<br>Learning rate scheduler<br>Gradient clipping value<br>Patience tolerance |

# Hyperparameter Optimization



# Automation and Reproducibility

# Automation and Reproducibility



# Data product: Strengths & Weaknesses

- 
- Not friendly to non-technical users
  - Not covering all methods
  - Implemented methods can only be used within the designed system
- Python package
  - High flexibility
  - Easy to make improvements
  - Generalized
  - Automated

Introduction & Objectives

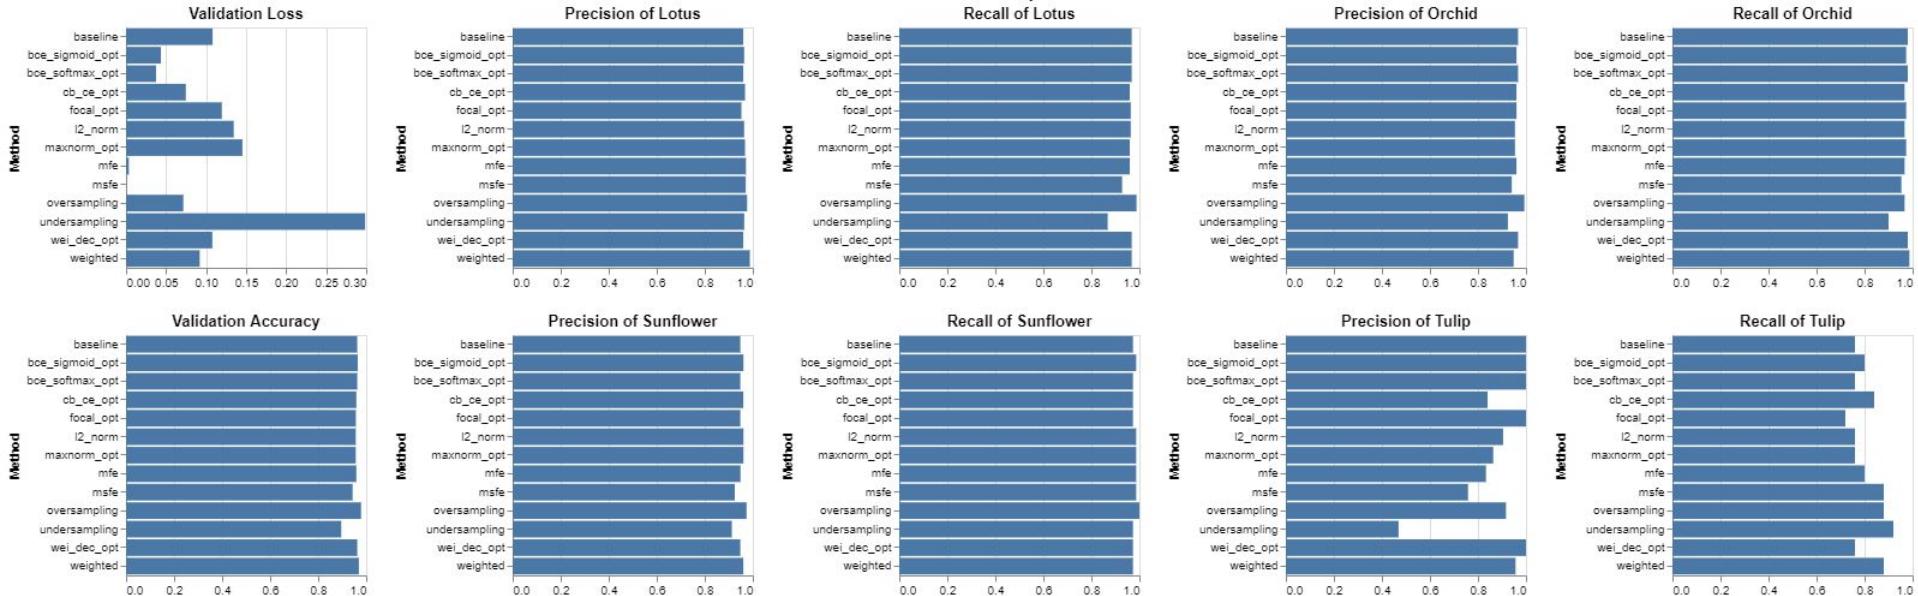
Data Science Techniques

Data Product

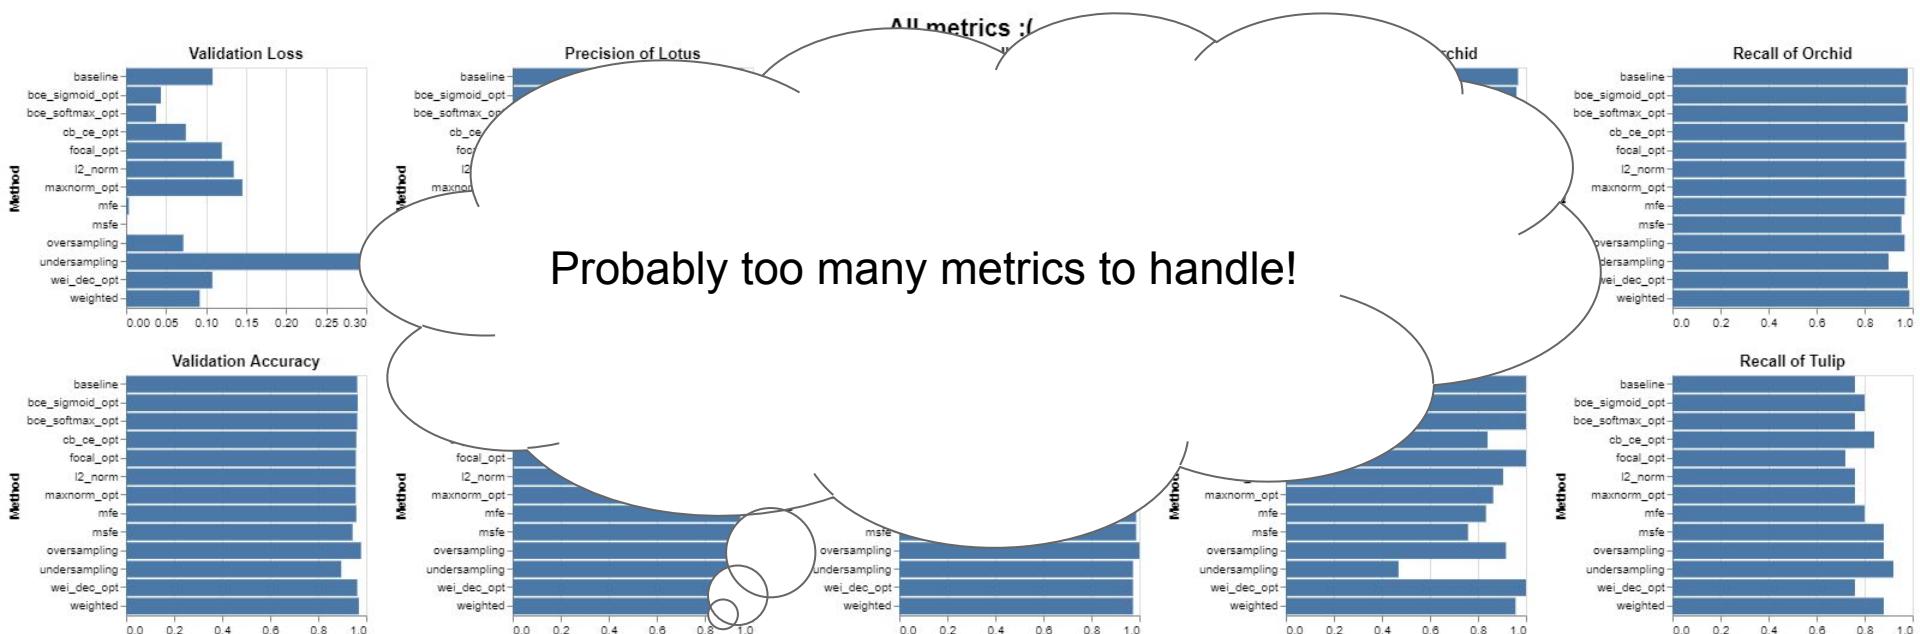
# Results

Conclusions and Recommendations

# Decision-making on “Best”



# Decision-making on “Best”



# Traditional Weighted Metrics



| Precision | <i>n</i> |
|-----------|----------|
| 0.7       | 30       |
| 0.8       | 20       |
| 0.3       | 10       |

$$\text{macro average} = \frac{0.7 + 0.8 + 0.3}{3} = 0.6$$

$$\text{weighted average} = \frac{0.7 \times 30 + 0.8 \times 20 + 0.3 \times 10}{30 + 20 + 10} = 0.6667$$

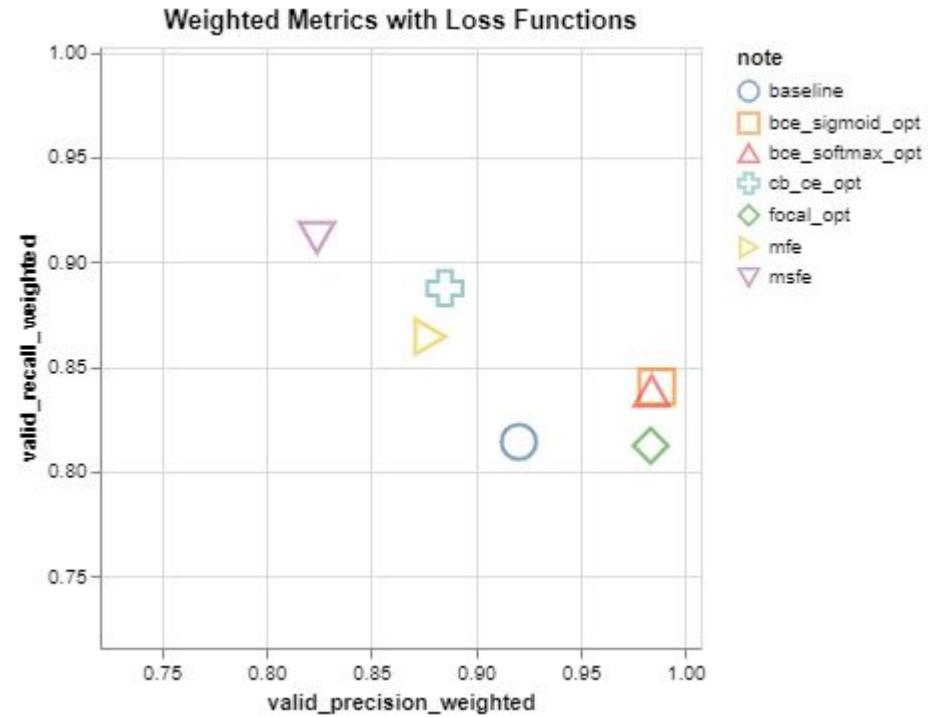
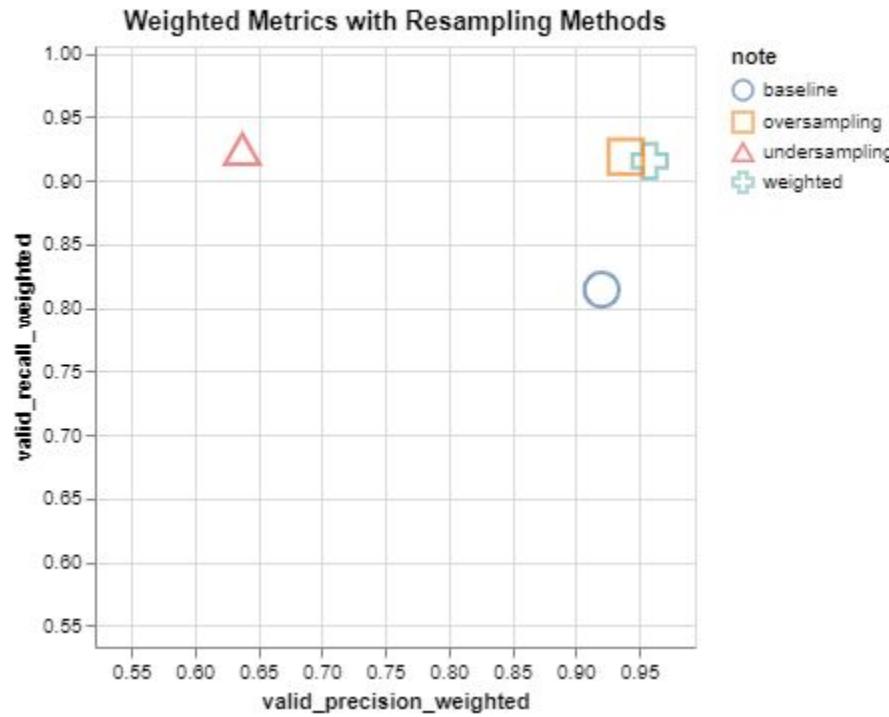
# Weighted Metrics



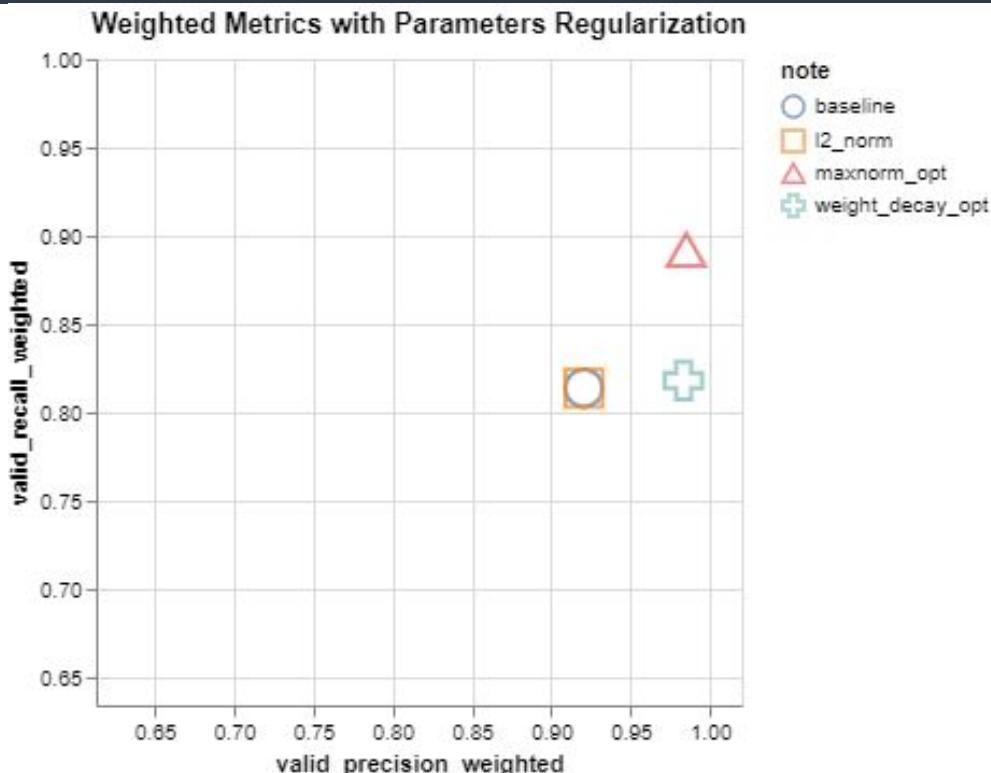
| Precision | $n$ | $1/n$ |
|-----------|-----|-------|
| 0.7       | 30  | 1/30  |
| 0.8       | 20  | 1/20  |
| 0.3       | 10  | 1/10  |

$$\text{new weighted average} = \frac{0.7 \times 1/30 + 0.8 \times 1/20 + 0.3 \times 1/10}{1/30 + 1/20 + 1/10} = 0.4$$

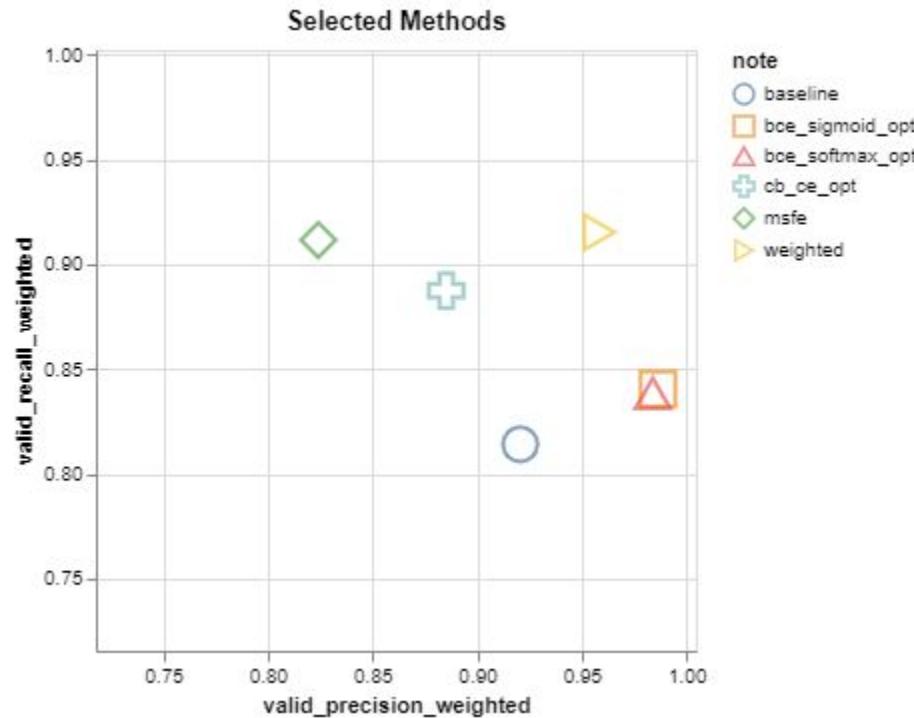
# Results



# Results (cont.)



# Results: Best Methods



Introduction & Objectives

Data Science Techniques

Data Product

Results

# **Conclusions and Recommendations**

# Conclusions and Recommendations

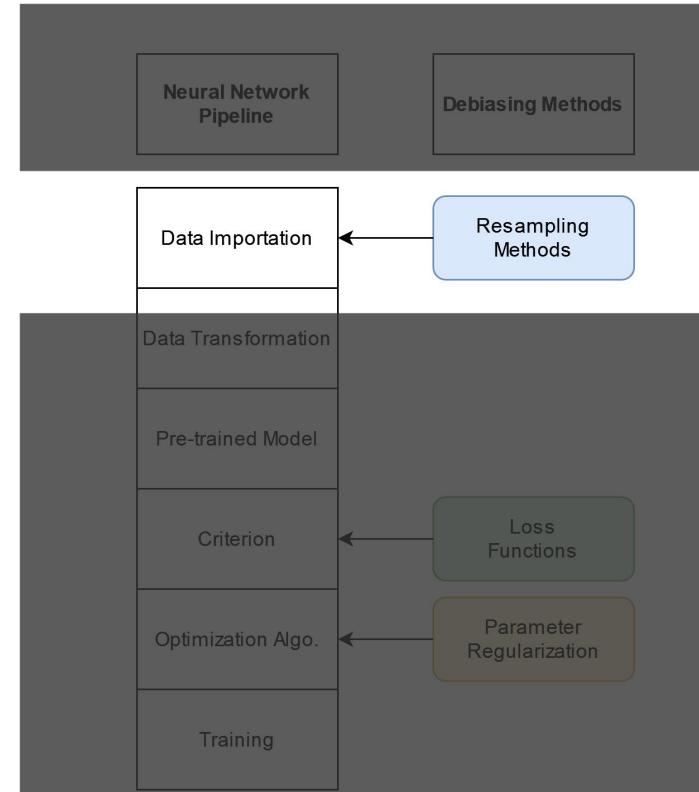
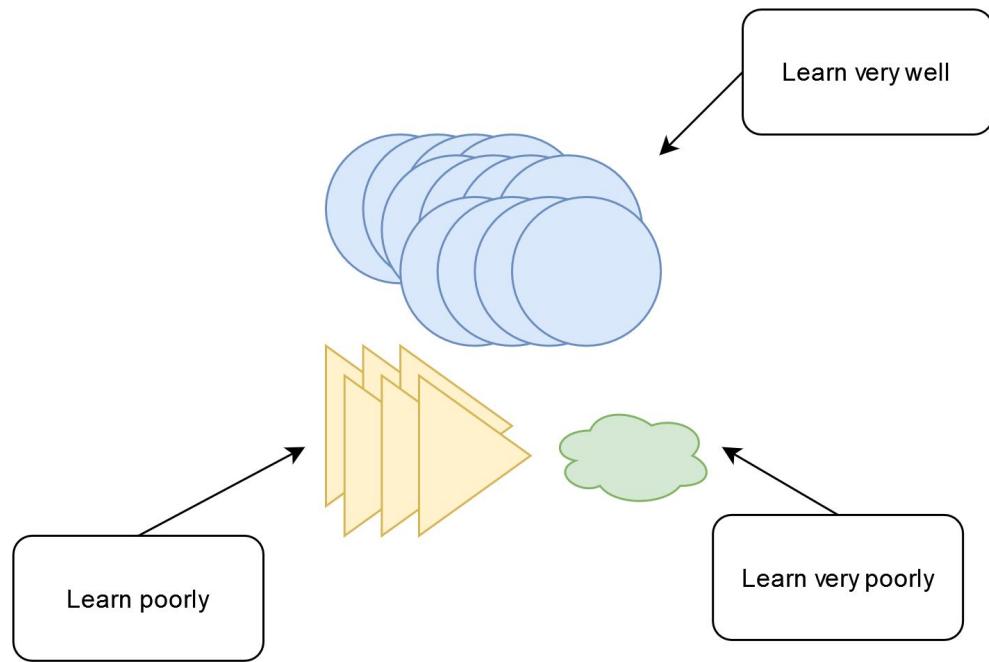
Created a package that is able to train and compare CNN models that can employ various techniques to tackle data imbalance problems

- ★ Successfully implemented various methods involving resampling techniques, parameter regularization, and custom loss functions to improve precision/recall and accuracy in the minority class
- ★ May be limited by data quality / data preprocessing steps
- ★ Add more features:
  - Adapting to more pytorch pretrained models
  - Exploring different weighing schemes for class-balanced loss
  - ROC curve analysis

# Thank you!

Any questions?

# Resampling



# Training

```
dict_transform = transformer(image_size=224,
 normalize=(0.5,0.5),
 vflip=True,
 hflip=True,
 rotate=True,
 degrees=20,
 crop=True)

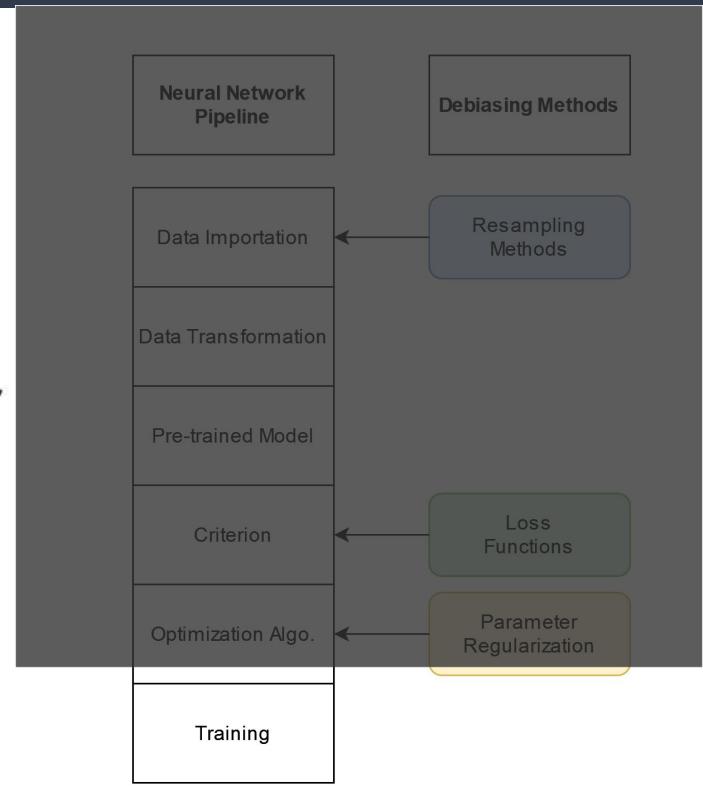
dict_model = load_model(model_id='resnet18',
 n_class=4,
 pretrained=True,
 full=False)

dict_dataloader = load_image(train_path='/data/train/',
 valid_path='/data/validation/',
 batch_size=64,
 dict_transform=dict_transform,
 sampling_strategy=None)

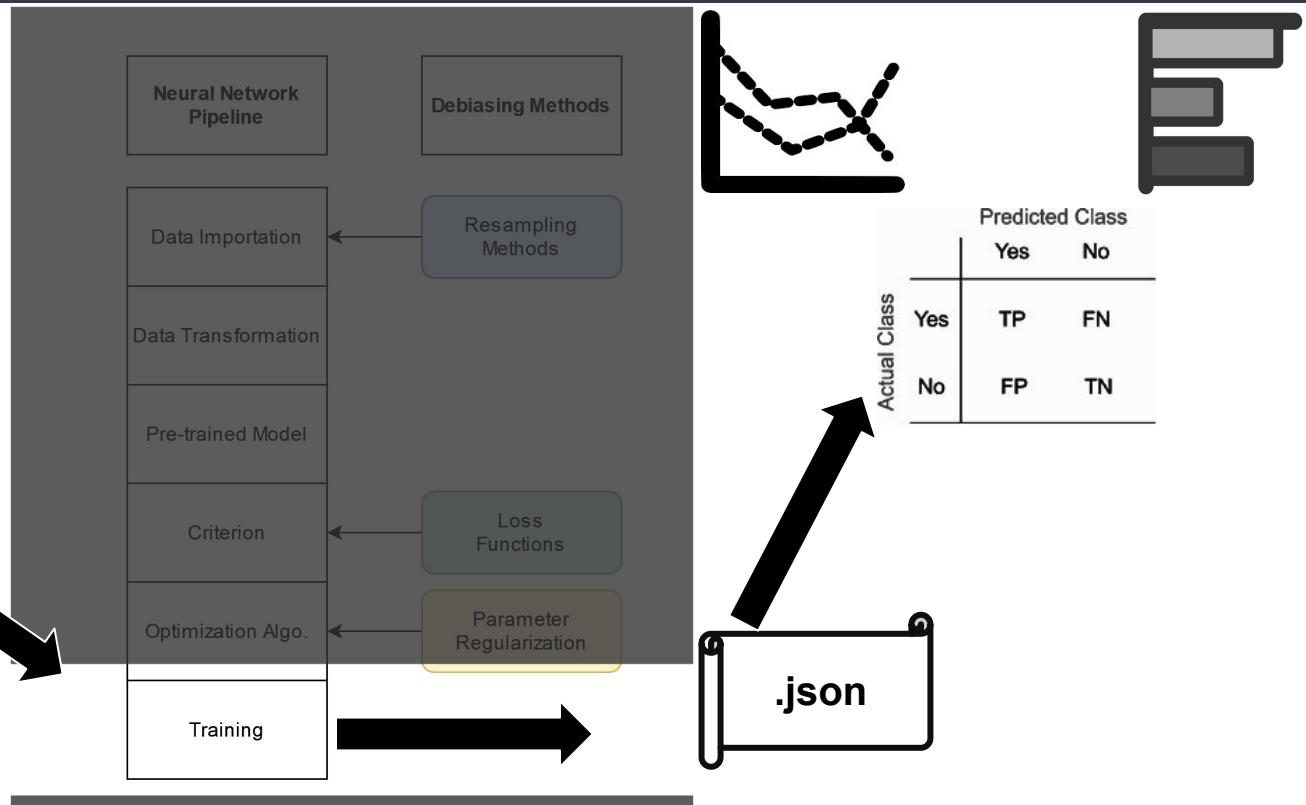
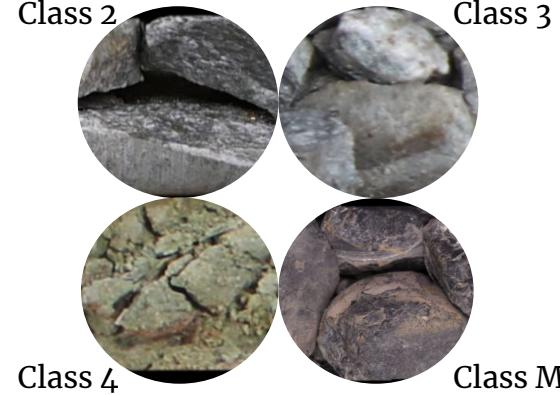
dict_optimizer = load_optimizer(dict_model=dict_model,
 lr=1e-6,
 optimizer_name='SGD',
 strategy=None,
 weight_decay=0.001,
 thresh=1.0)

dict_loss = load_loss(dict_dataloader=dict_dataloader,
 loss_family='CB',
 loss_type='CE',
 beta=0.999,
 gamma=None)

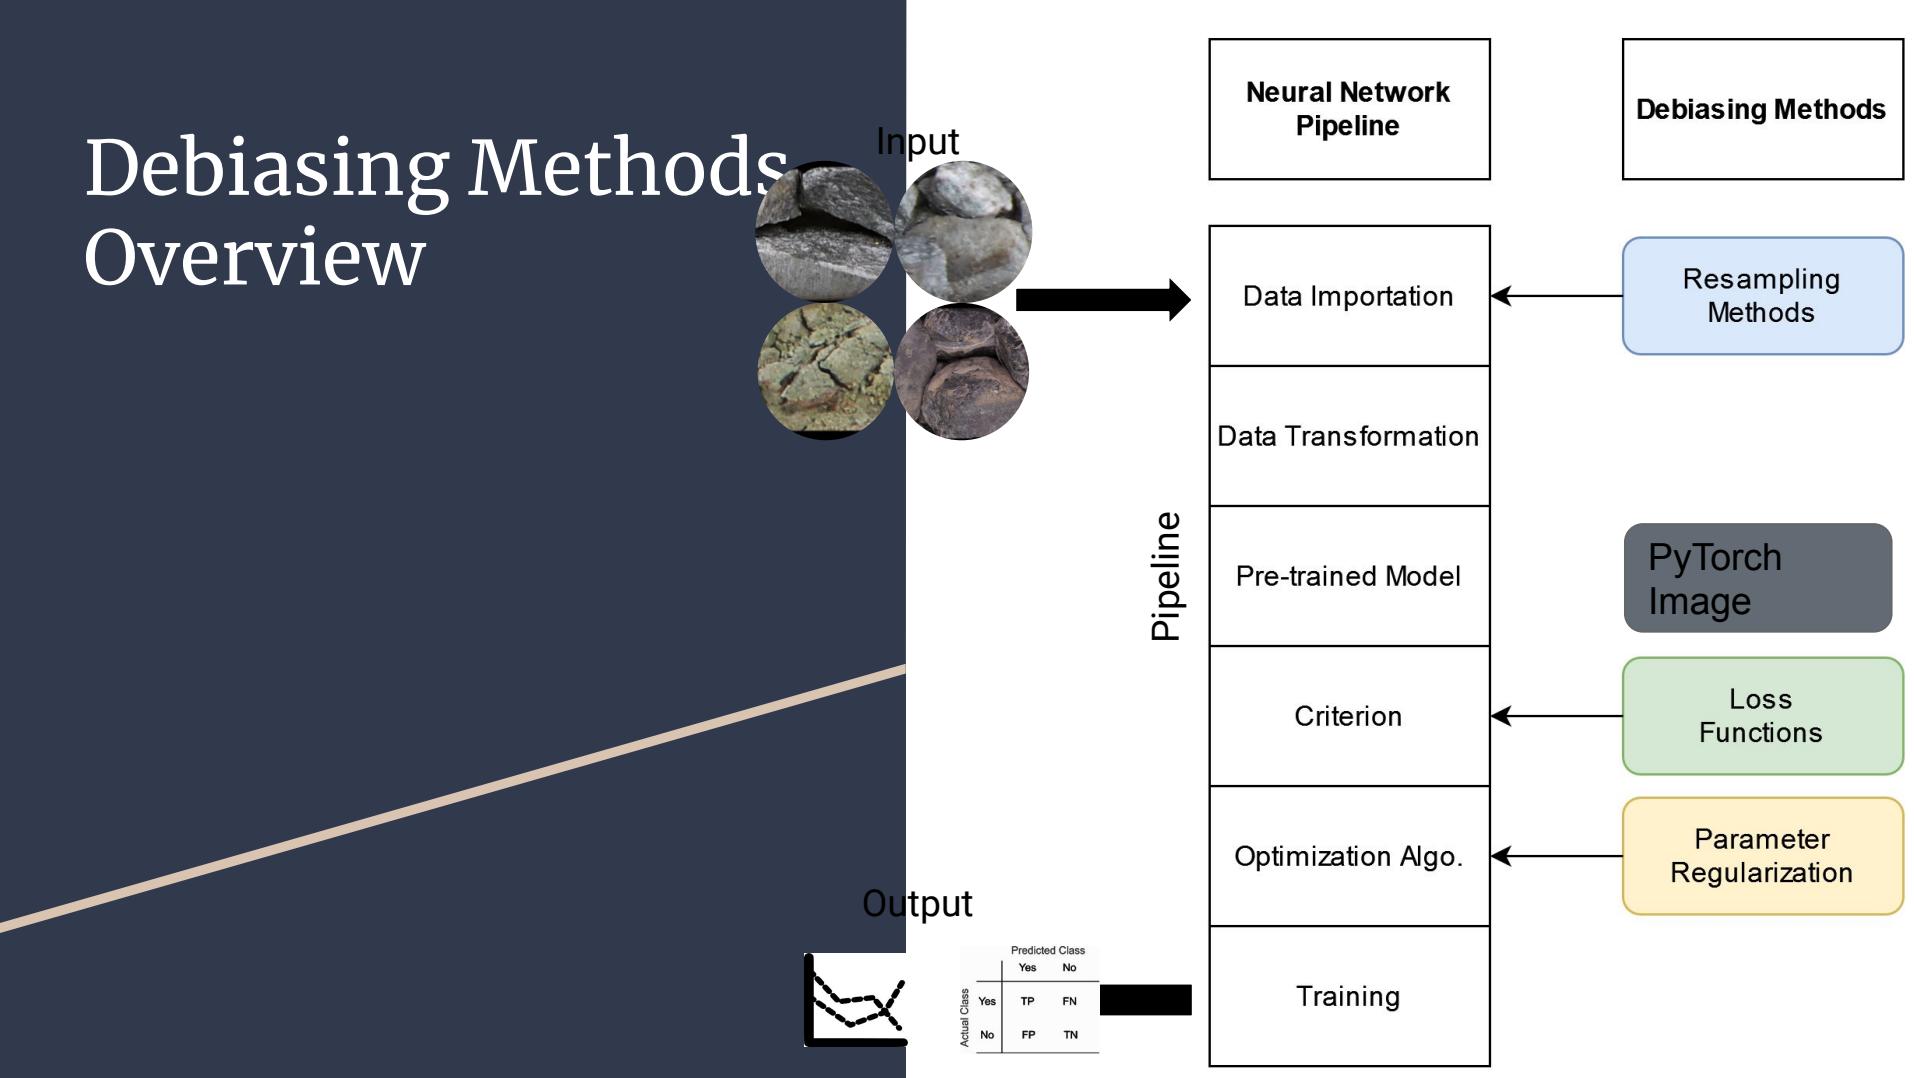
results = train(dict_model=dict_model,
 dict_loss=dict_loss,
 dict_optimizer=dict_optimizer,
 dict_dataloader=dict_dataloader,
 epochs=20,
 verbose=False,
 note='Testing - ResNet',
 log_id='Resnet-Rocks-lr=1e-6')
```



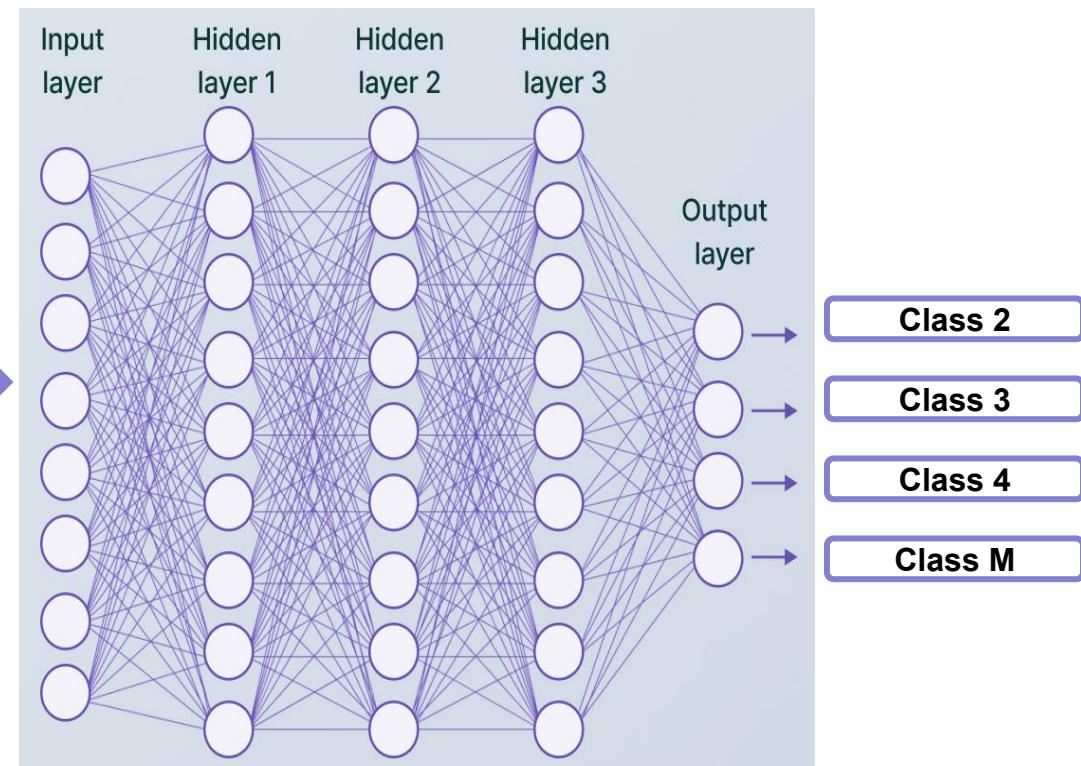
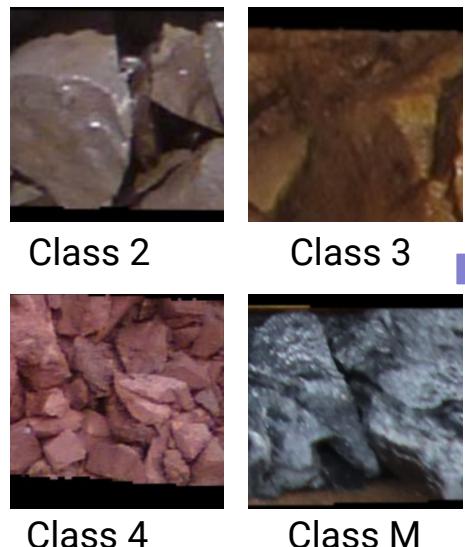
# Training



# Debiasing Methods Overview



# General CNN Structure



# ResNet18

| layer name | output size      | 18-layer                                                                                                                                    |
|------------|------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| conv1      | $112 \times 112$ |                                                                                                                                             |
| conv2_x    | $56 \times 56$   | <br>Basic Block<br>$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$<br>Number of times to repeat the Basic Block   |
| conv3_x    | $28 \times 28$   | <br>Basic Block<br>$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$<br>Number of times to repeat the Basic Block |
| conv4_x    | $14 \times 14$   | <br>Basic Block<br>$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$<br>Number of times to repeat the Basic Block |
| conv5_x    | $7 \times 7$     | <br>Basic Block<br>$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$<br>Number of times to repeat the Basic Block |
|            | $1 \times 1$     |                                                                                                                                             |
| FLOPs      |                  | $1.8 \times 10^9$                                                                                                                           |

Changed FC Layer  
from:

(fc): Linear(in\_features=512, out\_features=1000, bias=True)

to:

```
new_layers = nn.Sequential(
 nn.ReLU(),
 nn.Dropout(0.5),
 nn.Linear(512, n_class)
)
```

# MobileNet v2

| Input                    | Operator    | t | c    | n | s |
|--------------------------|-------------|---|------|---|---|
| $224^2 \times 3$         | conv2d      | - | 32   | 1 | 2 |
| $112^2 \times 32$        | bottleneck  | 1 | 16   | 1 | 1 |
| $112^2 \times 16$        | bottleneck  | 6 | 24   | 2 | 2 |
| $56^2 \times 24$         | bottleneck  | 6 | 32   | 3 | 2 |
| $28^2 \times 32$         | bottleneck  | 6 | 64   | 4 | 2 |
| $14^2 \times 64$         | bottleneck  | 6 | 96   | 3 | 1 |
| $14^2 \times 96$         | bottleneck  | 6 | 160  | 3 | 2 |
| $7^2 \times 160$         | bottleneck  | 6 | 320  | 1 | 1 |
| $7^2 \times 320$         | conv2d 1x1  | - | 1280 | 1 | 1 |
| $7^2 \times 1280$        | avgpool 7x7 | - | -    | 1 | - |
| $1 \times 1 \times 1280$ | conv2d 1x1  | - | k    | - | - |

Changed classifier  
from:

```
(classifier): Sequential(
 (0): Dropout(p=0.2, inplace=False)
 (1): Linear(in_features=1280, out_features=1000, bias=True)
)
```

to:

```
new_layers = nn.Sequential(
 nn.ReLU(),
 nn.Dropout(0.5),
 nn.Linear(1280, n_class)
)
```

# EfficientNet Bo

| Stage<br>$i$ | Operator<br>$\hat{\mathcal{F}}_i$ | Resolution<br>$\hat{H}_i \times \hat{W}_i$ | #Channels<br>$\hat{C}_i$ | #Layers<br>$\hat{L}_i$ |
|--------------|-----------------------------------|--------------------------------------------|--------------------------|------------------------|
| 1            | Conv3x3                           | $224 \times 224$                           | 32                       | 1                      |
| 2            | MBConv1, k3x3                     | $112 \times 112$                           | 16                       | 1                      |
| 3            | MBConv6, k3x3                     | $112 \times 112$                           | 24                       | 2                      |
| 4            | MBConv6, k5x5                     | $56 \times 56$                             | 40                       | 2                      |
| 5            | MBConv6, k3x3                     | $28 \times 28$                             | 80                       | 3                      |
| 6            | MBConv6, k5x5                     | $14 \times 14$                             | 112                      | 3                      |
| 7            | MBConv6, k5x5                     | $14 \times 14$                             | 192                      | 4                      |
| 8            | MBConv6, k3x3                     | $7 \times 7$                               | 320                      | 1                      |
| 9            | Conv1x1 & Pooling & FC            | $7 \times 7$                               | 1280                     | 1                      |

→ Changed classifier  
from:

```
(classifier): Sequential(
 (0): Dropout(p=0.2, inplace=True)
 (1): Linear(in_features=1280, out_features=1000, bias=True)
)
)
```

to:

```
new_layers = nn.Sequential(
 nn.ReLU(),
 nn.Dropout(0.5),
 nn.Linear(1280, n_class)
)
```

# JSON File Format

```
● ↴ {
 "note": "A custom note for the experimental run",
 "date": "20230609-134532",
 "results": {
 "train_loss": [0.9, 0.7, ..., 0.34],
 "train_accuracy": [0.8, ..., 0.99],
 "train_precision": {
 "ClassName1": [0.5, ..., 0.8],
 "ClassName2": [0.6, ..., 0.7]
 },
 "train_recall": {
 "ClassName1": [0.4, 0.7, ..., 0.7],
 "ClassName2": [0.8, 0.78, ..., 0.81]
 },
 "valid_loss": [0.9, 0.7, ..., 0.34],
 "valid_accuracy": [0.8, ..., 0.99],
 "valid_precision": {
 "ClassName1": [0.5, ..., 0.8],
 "ClassName2": [0.6, ..., 0.7]
 },
 "valid_recall": {
 "ClassName1": [0.4, 0.7, ..., 0.7],
 "ClassName2": [0.8, 0.78, ..., 0.81],
 }
 }
}
```

# JSON File Format

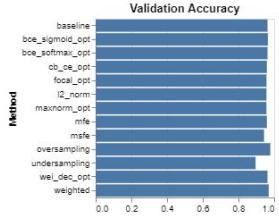
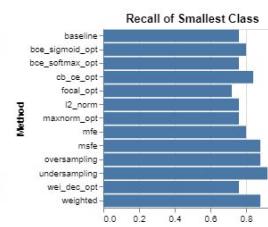
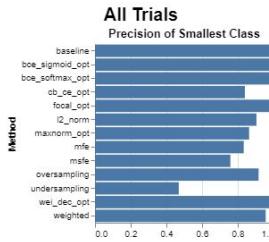
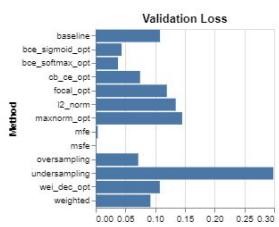
```
 },
 "y": {
 "train": {
 "true": {
 "0": [1, 0, 1, 0, ..., 0], # Epoch 0
 "1": [0, 1, 1, 0, ..., 0], # Epoch 1...
 },
 "pred": {
 "0": [0, 0, 1, 0, ..., 0], # Epoch 0
 "1": [1, 0, 0, 0, ..., 0], # Epoch 1...
 }
 },
 "valid": {
 "true": {
 "0": [0, 1, 0, 1, ..., 1], # Epoch 0
 "1": [0, 1, 1, 0, ..., 0], # Epoch 1...
 },
 "pred": {
 "0": [0, 0, 1, 0, ..., 1], # Epoch 0
 "1": [0, 1, 1, 1, ..., 0], # Epoch 1...
 }
 }
 },
 "var_dict": {
```

# JSON File Format

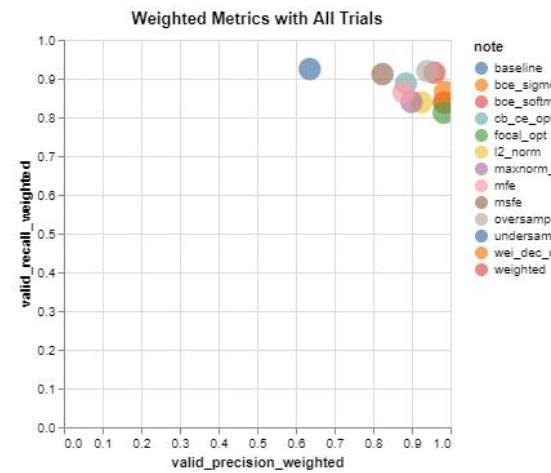
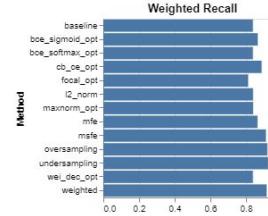
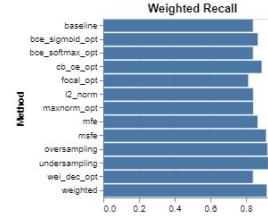
```
 },
 "var_dict": {
 "batch_size": 32,
 "sampling_strategy": None,
 "image_size": 224.0,
 "normalize_mean": 0.5,
 "normalize_std": 0.5,
 "vflip": False,
 "hflip": False,
 "rotate": True,
 "degrees": 30.0,
 "model_name": "resnet18",
 "n_class": 4,
 "loss_family": "CB",
 "loss_type": "focal",
 "beta": 0.999,
 "gamma": 0.5,
 "optimizer": "Adam",
 "reg_strategy": None,
 "learning_rate": 1e-05,
 "epochs": 40,
 "patience": 5
 },
 }
```

# Statistics and Visualization

|   | valid_loss | valid_accuracy | sampling_strategy | loss_family | loss_type        | reg_strategy | learning_rate | note             |
|---|------------|----------------|-------------------|-------------|------------------|--------------|---------------|------------------|
| 2 | 0.071542   | 0.978516       | oversampling      | baseline    | None             | None         | 1.00e-05      | oversampling_opt |
| 1 | 0.091995   | 0.969531       | weighted          | baseline    | None             | None         | 1.00e-05      | weighted_opt     |
| 6 | 0.043558   | 0.964844       | None              | CB          | BCE with sigmoid | None         | 1.00e-05      | bce_sigmoid_opt  |
| 0 | 0.108101   | 0.962891       | None              | baseline    | None             | None         | 1.00e-05      | baseline_opt     |



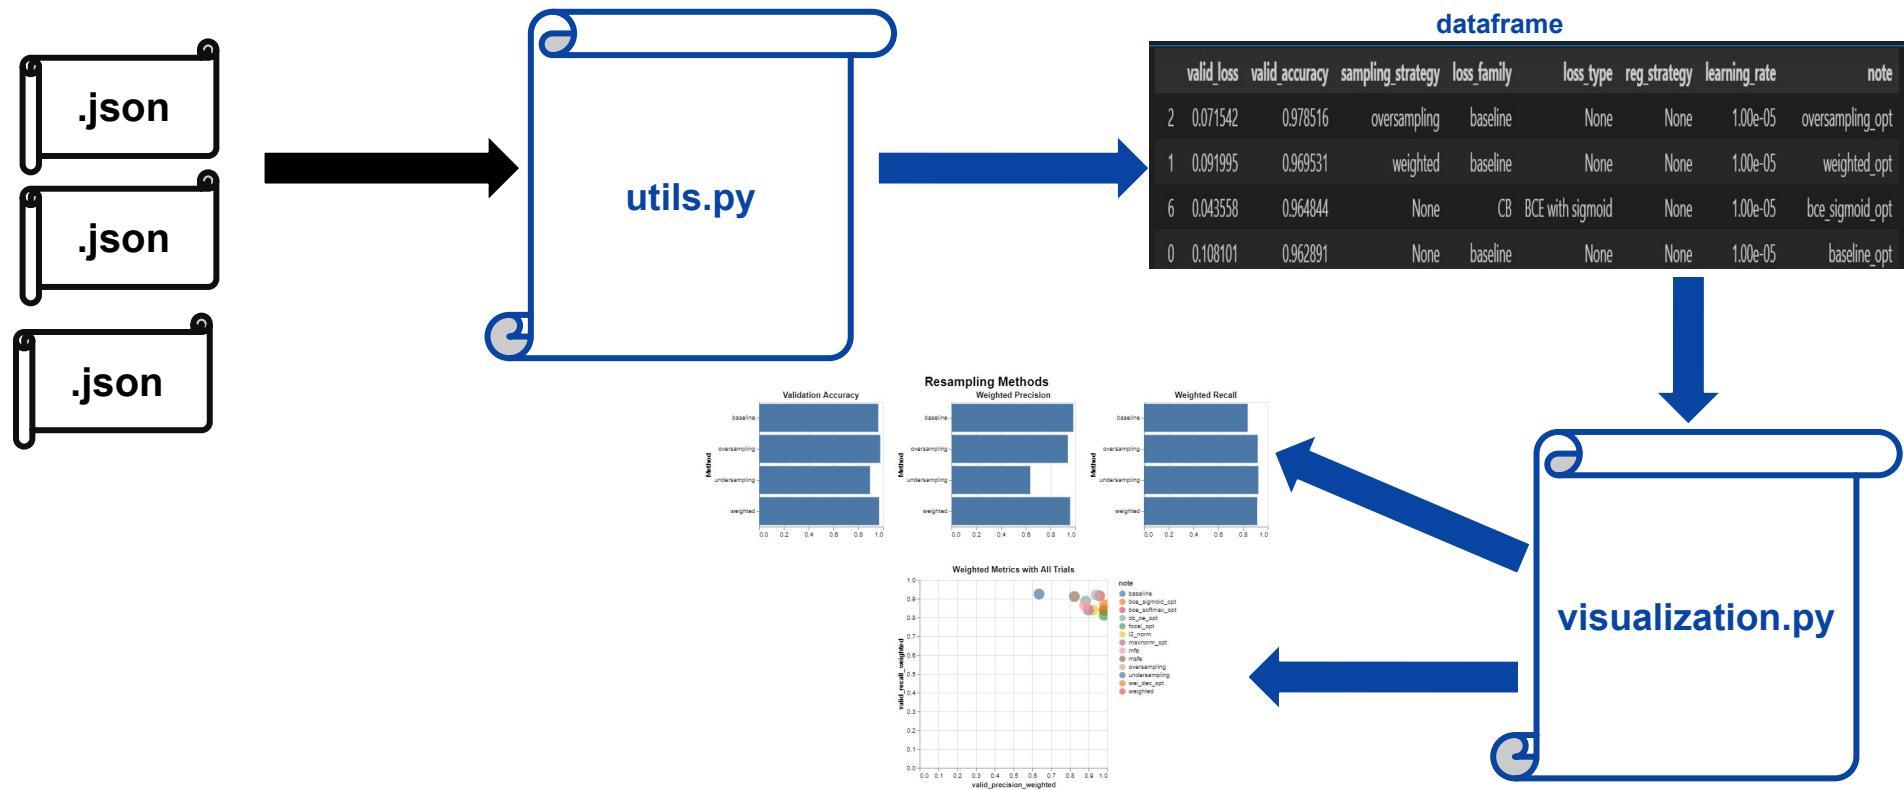
Method



note

- baseline
- bce\_sigmoid\_opt
- bce\_softmax\_opt
- cb\_ce\_opt
- focal\_opt
- l2\_norm
- maxnorm\_opt
- mfe
- msfe
- oversampling
- undersampling
- wei\_deo\_opt
- weighted

# Graphing Multiple JSON files



# Loss Functions – Effective Number

$$E_n = (1 - \beta^n) / (1 - \beta), \text{ where } \beta \in [0, 1]:$$

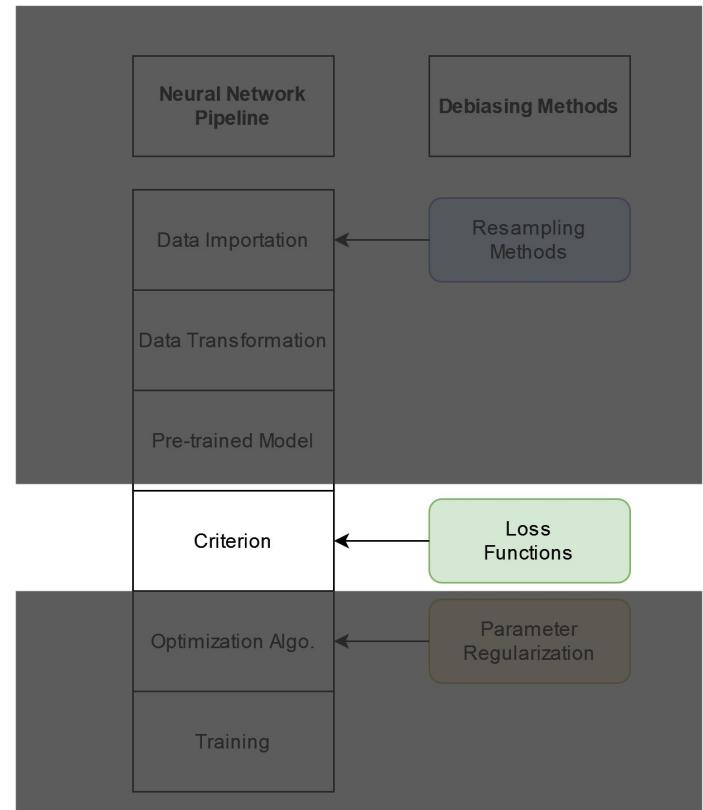
- $\beta = 0$  : no re-weighing
- $\beta = 1$  : re-weighing by inverse class frequencies

So, class-balanced (CB) loss will be defined as:

$$\text{CB}(\mathbf{p}, y) = \frac{1}{E_{n_y}} \mathcal{L}(\mathbf{p}, y) = \frac{1 - \beta}{1 - \beta^{n_y}} \mathcal{L}(\mathbf{p}, y)$$

Where  $\mathbf{p}$  is class probabilities,  $y$  is ground-truth labels,  $n_y$  is the number of samples in the ground-truth class and  $\mathcal{L}(\mathbf{p}, y)$  is any loss function.

(ref)



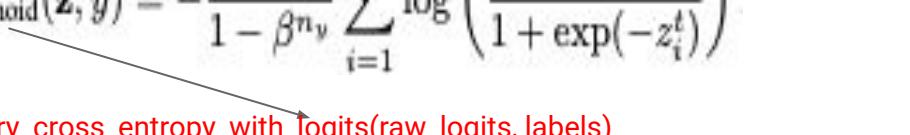
# Loss Functions - Class-Balanced Loss

$$\text{CB}_{\text{softmax}}(\mathbf{z}, y) = -\frac{1-\beta}{1-\beta^{n_y}} \log \left( \frac{\exp(z_y)}{\sum_{j=1}^C \exp(z_j)} \right)$$

cross\_entropy(raw\_logits, labels)      binary\_crossentropy(softmax\_output, labels)

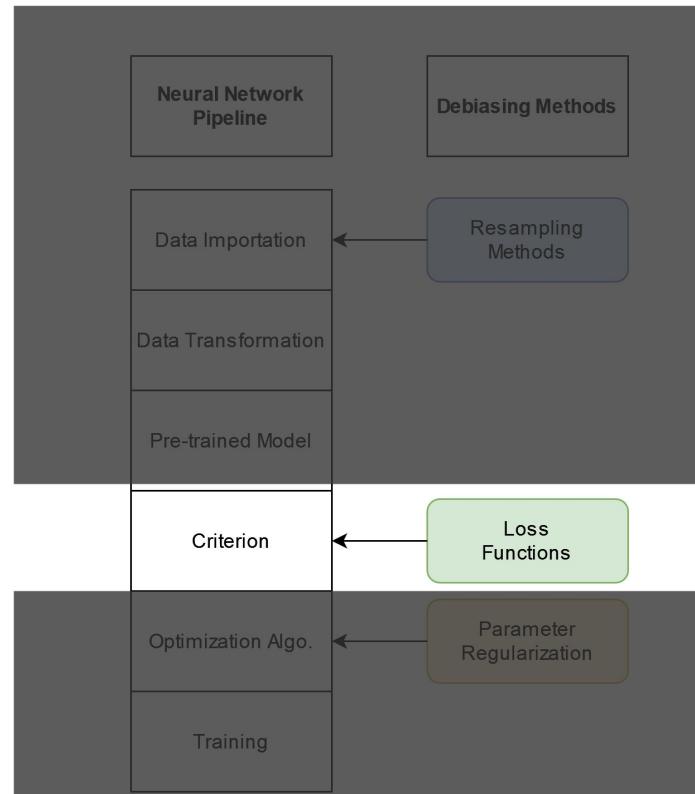
$$\text{CB}_{\text{sigmoid}}(\mathbf{z}, y) = -\frac{1-\beta}{1-\beta^{n_y}} \sum_{i=1}^C \log \left( \frac{1}{1+\exp(-z_i^t)} \right)$$

binary\_cross\_entropy\_with\_logits(raw\_logits, labels)



$$\text{CB}_{\text{focal}}(\mathbf{z}, y) = -\frac{1-\beta}{1-\beta^{n_y}} \sum_{i=1}^C (1-p_i^t)^\gamma \log(p_i^t).$$

where  $y$  is called the modulating factor



# Parameter Regularization

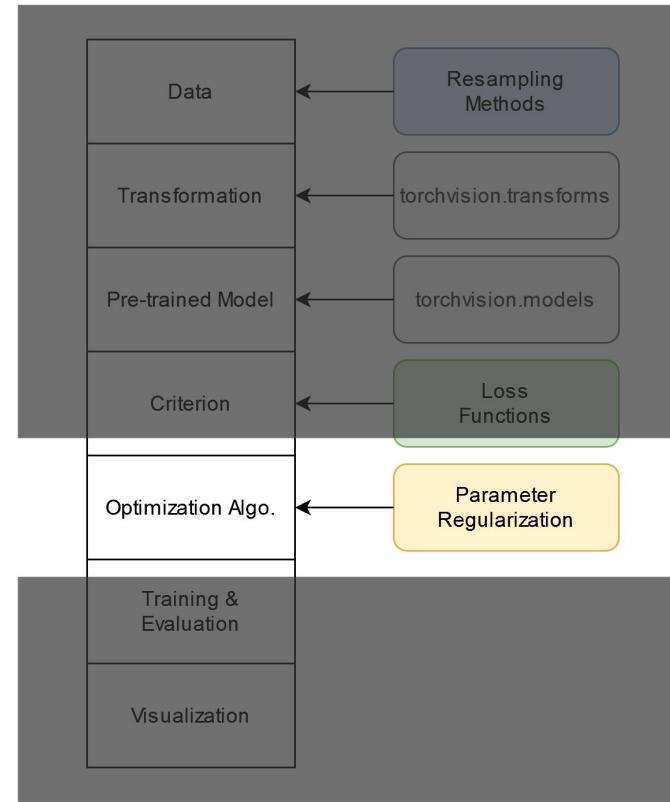
Weight balancing:

- A naive model has “artificially” large weights for common classes
- Try to balance the weights across classes for better predictions on the minority class

$$\Theta^* = \arg \min_{\Theta} F(\Theta; \mathcal{D}) \equiv \sum_{i=1}^N \ell(f(\mathbf{x}_i; \Theta), y_i)$$

Optimized set of weights per layer of the network  
[\(ref\)](#)

Neural network      Loss function



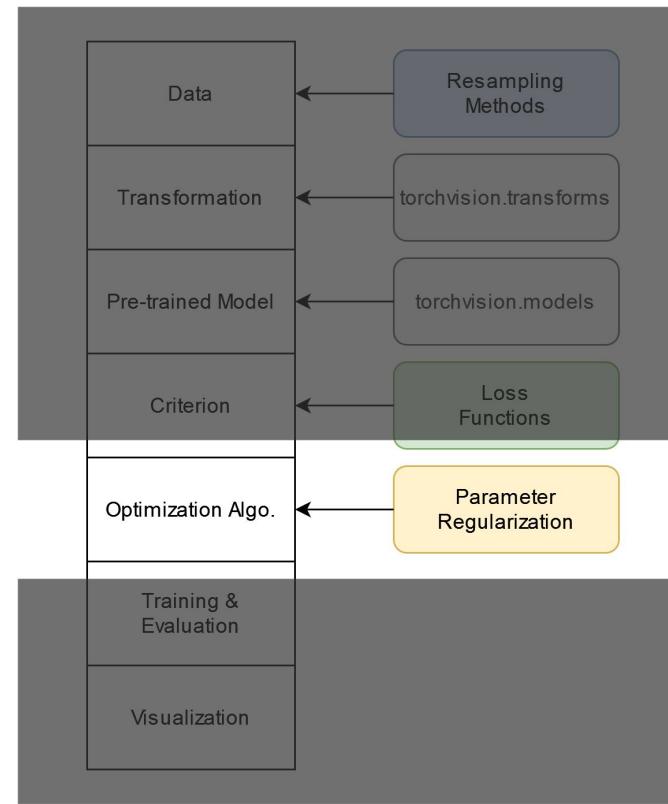
# Parameter Regularization

$$\boldsymbol{\Theta}^* = \arg \min_{\boldsymbol{\Theta}} F(\boldsymbol{\Theta}; \mathcal{D}) + \lambda \sum_k \|\boldsymbol{\theta}_k\|_2^2,$$

**Hyperparameter**      **Weights for each class  $k$**

Weight decay limits weight growth and penalizes larger weights more severely than smaller weights, which promotes more balanced weight growth.

Weight decay decreases model complexity, which mitigates overfitting and improves generalization



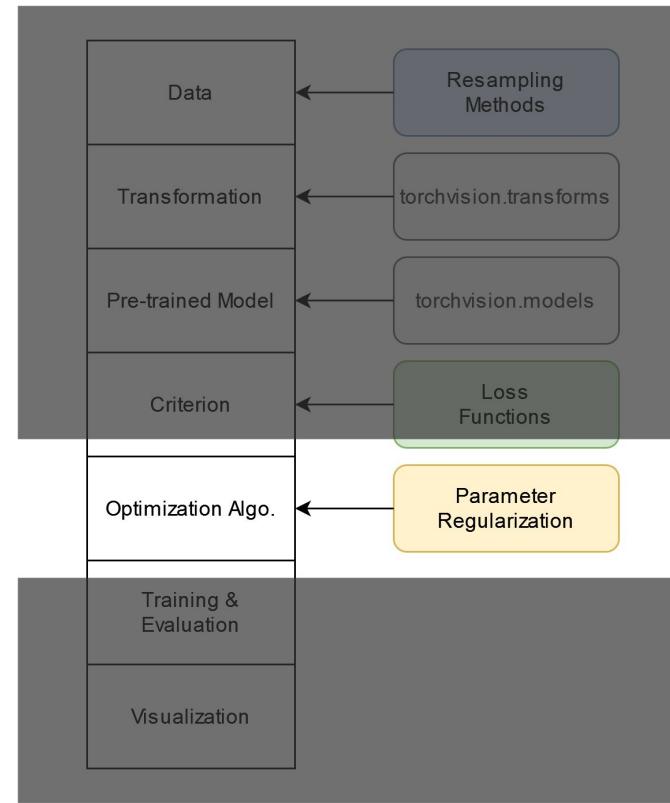
# Parameter Regularization

L2- normalization perfectly balances classifier weights to be the unit norm

$$\Theta^* = \arg \min_{\Theta} F(\Theta; \mathcal{D}), \quad s.t. \quad \|\theta_k\|_2^2 = 1, \quad \forall k.$$

Weights for each class  $k$

Improves performance on the minority classes, but sacrifices common-class accuracy and overall model performance.



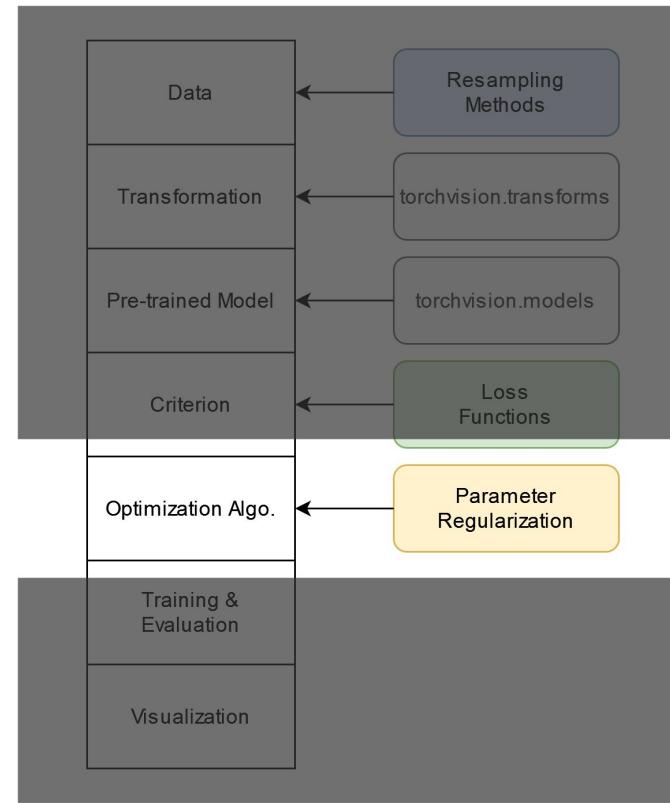
# Parameter Regularization

MaxNorm encourages smaller weight growth and places a limit on weight norms ( $\delta$ )

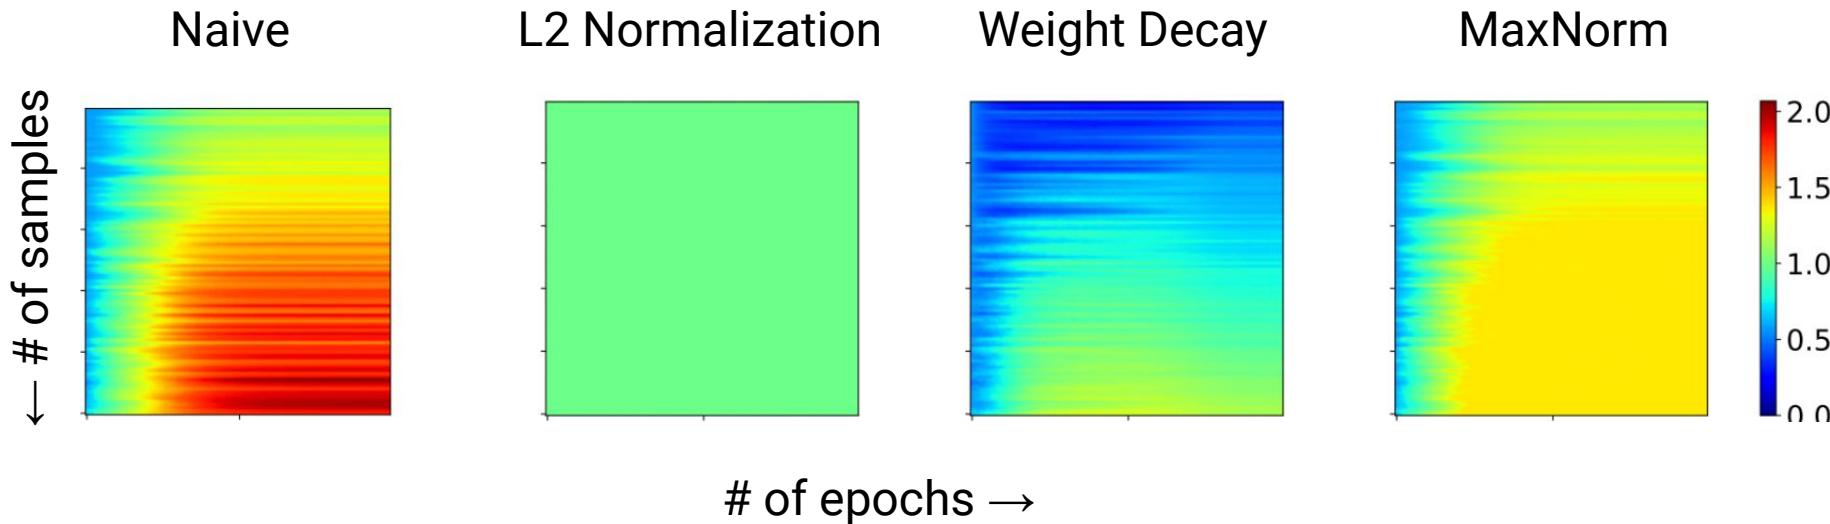
$$\Theta^* = \arg \min_{\Theta} F(\Theta; \mathcal{D}), \quad s.t. \quad \|\theta_k\|_2^2 \leq \delta^2, \quad \forall k,$$

Weights for each class  $k$

Weights cannot “explode” even at high learning rates because they are always bounded.



# Parameter Regularization



(Source)

# Hyperparameter Optimization

Base Debiased Functions

```
transformer()
• Image
 transformations
• Stored in
 dict_transform
load_model()
• Model selection
• Stored in
 dict_model
```

Parameters to Tune

```
params = {
 'optimizer_strategy': ['maxnorm'],
 'optimizer_thresh': [0.5, 2],
 'criterion_loss_family': ["CB"],
 'criterion_beta': [0.5, 0.9],
 'criterion_loss_type': ["CE"]}
```

Main Tuning Function

```
manual_opt():
 file_path,
 dict_transform,
 dict_model,
 params,
 best_by
 cv,
 sampling_strategy,
 random_state
```

# Hyperparameter Optimization

Output:

- All parameter combinations and their associated **mean validation accuracy**, and **mean precision and recall** for the minority class
- The metric used to determine the “best” cross-validation result
- The **best** parameter combination and its associated metric

